

Query Optimization: From 0 to 10 (and up to 5.7)

Jaime Crespo

Percona Live Europe 2015

-Amsterdam, 21 Sep 2015-

dbahire.com/pleu15

Agenda - First 3 hours

1. Introduction	5. FULLTEXT Search
2. Access Types and Basic Indexing Techniques	6. Joins
3. Break	7. Subqueries
4. Multi-Column Indexing	8. Query Profiling

Agenda - Last 3 hours

1. General Optimizer Improvements	6. SQL Mode Changes
2. Computed/Virtual Columns	7. GIS Improvements and JSON Type
3. Query Rewrite Plugins	8. Results and Conclusions
4. Break	9. Q&A
5. Optimizer Hints	

Query Optimization: From 0 to 10 (and up to 5.7)

INTRODUCTION

This is me fighting bad query performance



- Sr. Database Administrator at Wikimedia Foundation
- Used to work as a trainer for Oracle (MySQL), as a Consultant (Percona) and as a Freelance administrator (DBAHire.com)

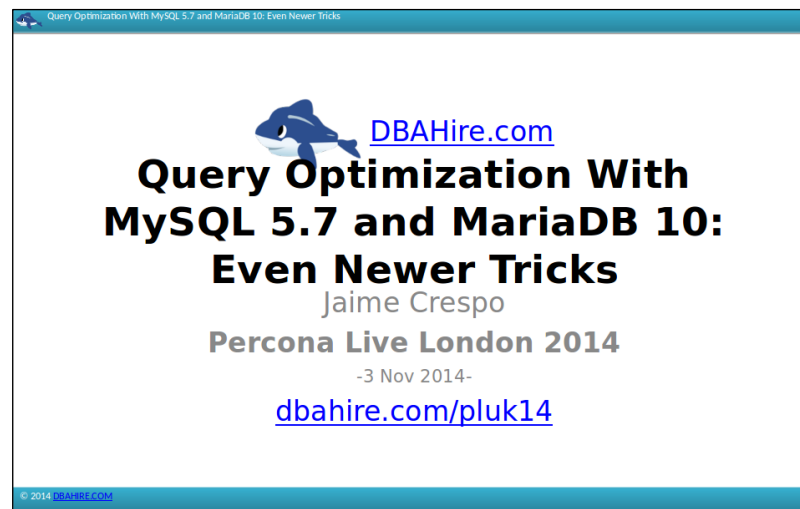
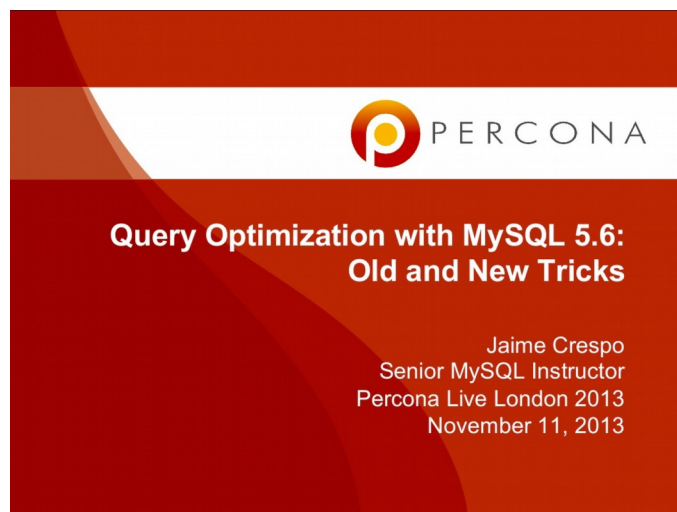
MySQL Versions

- 5.1 no longer has official support
- I will be showing you the results on mysql versions 5.5-5.7/10.1
- MySQL 5.7 and MariaDB 10.1 in RC with great new features

Recently Added Features Related to Query Optimization

- Index Condition Pushdown
- Subquery Optimizations (materialization and semijoin)
- IN-to-EXISTS/EXISTS-to-IN
- JOIN-to-WHERE
- Multi-Range Read
- Batched Key Access
- Persistent InnoDB Statistics
- UNION ALL optimization
- Improved GIS support
- EXPLAIN FORMAT=JSON
- EXPLAIN INSERT/UPDATE/DELETE
- Hash Join
- New optimizer hints
- New cost-based optimizer
- Optimizer Trace
- Filesort optimizations
- Virtual/computed columns and “functional indexes”
- New JSON type

I Already Mentioned Some of Those Topics during the Last Years



- Check my presentations here:
<http://www.slideshare.net/jynus/>

Example Application (I)

The screenshot shows the Wiktionary page for the word 'Anestesia'. The page layout includes a left sidebar with navigation links, a top navigation bar, and a main content area. The sidebar contains links such as 'Main page', 'Recent changes', 'New pages', 'Willekeurig woord', 'randompage-nld', 'prefixindex2', 'Categories', 'Informatie', 'De kroeg', 'Help', 'helpdesk', 'Financieel bijdragen', 'Zusterprojecten', 'Wikibooks', 'Wikipedia', 'Wikiquote', 'Wikisource', 'Commons', 'Tools', 'Links naar deze pagina', 'Related changes', 'Special pages', and 'Printable version'. The top navigation bar includes 'Page', 'Discussion', 'Read', 'Edit', 'View history', and a search box. The main content area features a 'Contents' table of contents, followed by the 'Italiaans' section with a category 'Woorden in het Italiaans' and a definition for 'Zelfstandig naamwoord'. The 'Spaans' section is partially visible at the bottom.

- Wiktionary (and all Wikimedia project's data) is licensed under the Creative Commons BY-SA-2.5 License and is Copyright its Contributors

Example Application (II)



- OSM Database is licensed under the Open DataBase License and is Copyright OpenStreetMap Contributors

Install the example databases

- Downloads and instructions at:
<http://dbahire.com/pleu15>
 - Requirements: a MySQL or MariaDB installation (MySQL Sandbox is suggested)
 - The wiktionary and OSM extracts
- Import them by doing:
`$ bzcat <file> | mysql <database>`

Query Optimization: From 0 to 10 (and up to 5.7)

ACCESS TYPES AND BASIC INDEXING TECHNIQUES

EXPLAIN

- Essential to understand the execution plan of our queries
 - Works on SELECTs, INSERTs, UPDATEs, REPLACEs, DELETEs and connections
 - Fully documented on:
<http://dev.mysql.com/doc/refman/5.6/en/explain-output.html>

EXPLAIN Example

MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page WHERE page_title = 'Dutch';

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	page	ALL	NULL	NULL	NULL	90956	Using where	

1 row in set (0.00 sec)

Difficult to see something

EXPLAIN Example (vertical format)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page WHERE  
page_title = 'Dutch'\G
```

```
***** 1. row *****  
      id: 1  
  select_type: SIMPLE  
        table: page  
         type: ALL  
possible_keys: NULL  
         key: NULL  
      key_len: NULL  
         ref: NULL  
        rows: 90956  
   Extra: Using where  
1 row in set (0.00 sec)
```



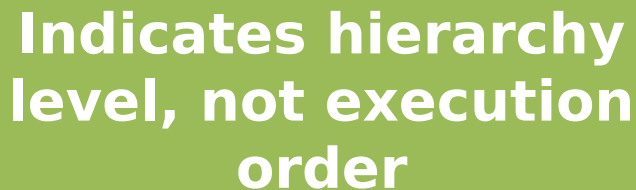
**Use \G for
vertical
formatting**

EXPLAIN Example (id)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page WHERE  
page_title = 'Dutch'\G
```

```
***** 1. row *****
```

```
id: 1  
select_type: SIMPLE  
table: page  
type: ALL  
possible_keys: NULL  
key: NULL  
key_len: NULL  
ref: NULL  
rows: 90956  
Extra: Using where  
1 row in set (0.00 sec)
```



**Indicates hierarchy
level, not execution
order**

EXPLAIN Example (select_type)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page WHERE
page_title = 'Dutch'\G
```

```
***** 1. row *****
```


```
      id: 1
  select_type: SIMPLE
        table: page
         type: ALL
possible_keys: NULL
          key: NULL
       key_len: NULL
          ref: NULL
         rows: 90956
      Extra: Using where
1 row in set (0.00 sec)
```

**Not a subquery or a
UNION**

EXPLAIN Example (table)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page WHERE  
page_title = 'Dutch'\G
```

```
***** 1. row *****  
      id: 1  
    select_type: SIMPLE  
      table: page  
        type: ALL  
possible_keys: NULL  
         key: NULL  
    key_len: NULL  
         ref: NULL  
        rows: 90956  
    Extra: Using where  
1 row in set (0.00 sec)
```



**Table scanned for
this step**

EXPLAIN Example (type)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page WHERE  
page_title = 'Dutch'\G
```

```
***** 1. row *****
```

```
    id: 1  
  select_type: SIMPLE  
    table: page  
    type: ALL  
possible_keys: NULL  
    key: NULL  
  key_len: NULL  
    ref: NULL  
   rows: 90956  
  Extra: Using where  
1 row in set (0.00 sec)
```

**All rows are read for
this table (FULL
TABLE SCAN)**

EXPLAIN Example (rows)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page WHERE  
page_title = 'Dutch'\G
```

```
***** 1. row *****  
      id: 1  
    select_type: SIMPLE  
      table: page  
       type: ALL  
possible_keys: NULL  
       key: NULL  
    key_len: NULL  
       ref: NULL  
     rows: 90956  
  Extra: Using where  
1 row in set (0.00 sec)
```

**Estimated number
of rows to be read
(all table rows)**

How to improve performance?

```
MariaDB [nlwiktionary]> SELECT * FROM page  
WHERE page_title = 'Dutch';  
1 row in set (0.11 sec)
```

- Let's add an index on page.page_title:

```
MariaDB [nlwiktionary]> ALTER TABLE page ADD INDEX  
page_title (page_title);  
Query OK, 0 rows affected (0.19 sec)  
Records: 0 Duplicates: 0 Warnings: 0
```

Index creation results (type)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page WHERE  
page_title = 'Dutch'\G
```

```
***** 1. row *****
```

```
    id: 1  
  select_type: SIMPLE  
    table: page  
    type: ref  
possible_keys: page_title  
    key: page_title  
  key_len: 257  
    ref: const  
    rows: 1  
  Extra: Using index condition  
1 row in set (0.00 sec)
```

type: ref means that an equality comparison will be checked against an index and several results could be returned

Index creation results (possible_keys and key)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page WHERE  
page_title = 'Dutch'\G
```

```
***** 1. row *****
```


```
    id: 1  
  select_type: SIMPLE  
    table: page  
    type: ref  
possible_keys: page_title  
    key: page_title  
  key_len: 257  
    ref: const  
   rows: 1  
  Extra: Using index condition  
1 row in set (0.00 sec)
```

**index(es) that the
optimizer considered
potentially useful, and
final index chosen**

Index creation results (ref)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page WHERE  
page_title = 'Dutch'\G
```

```
***** 1. row *****  
      id: 1  
    select_type: SIMPLE  
      table: page  
        type: ref  
possible_keys: page_title  
          key: page_title  
    key_len: 257  
      ref: const  
     rows: 1  
  Extra: Using index  
1 row in set (0.00 sec)
```



**Index is compared
with a constant, not
with another table**

Index creation results (rows)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page WHERE
page_title = 'Dutch'\G
```

```
***** 1. row *****
      id: 1
    select_type: SIMPLE
        table: page
         type: ref
possible_keys: page_title
         key: page_title
        key_len: 257
         ref: const
         rows: 1
    Extra: Using index
1 row in set (0.00 sec)
```

Only 1 row read. In this case, estimation is exact (thanks to index dive)

Index creation results (query time)

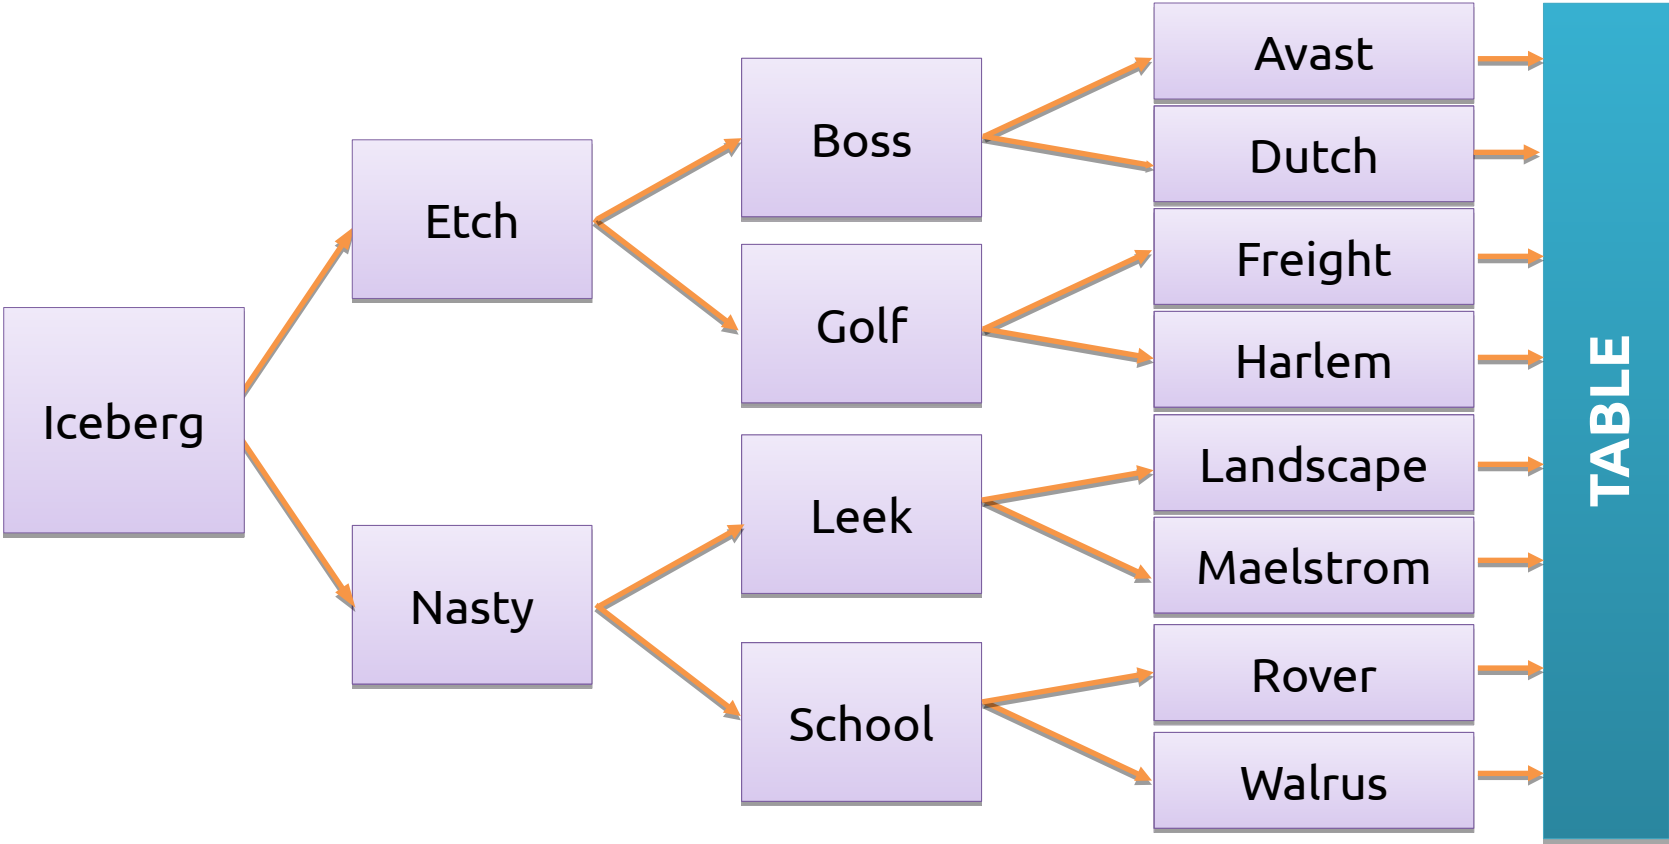
```
MariaDB [nlwiktionary]> SELECT * FROM page  
WHERE page_title = 'Dutch';  
1 row in set (0.00 sec)
```

**Query time has been
reduced substantially**

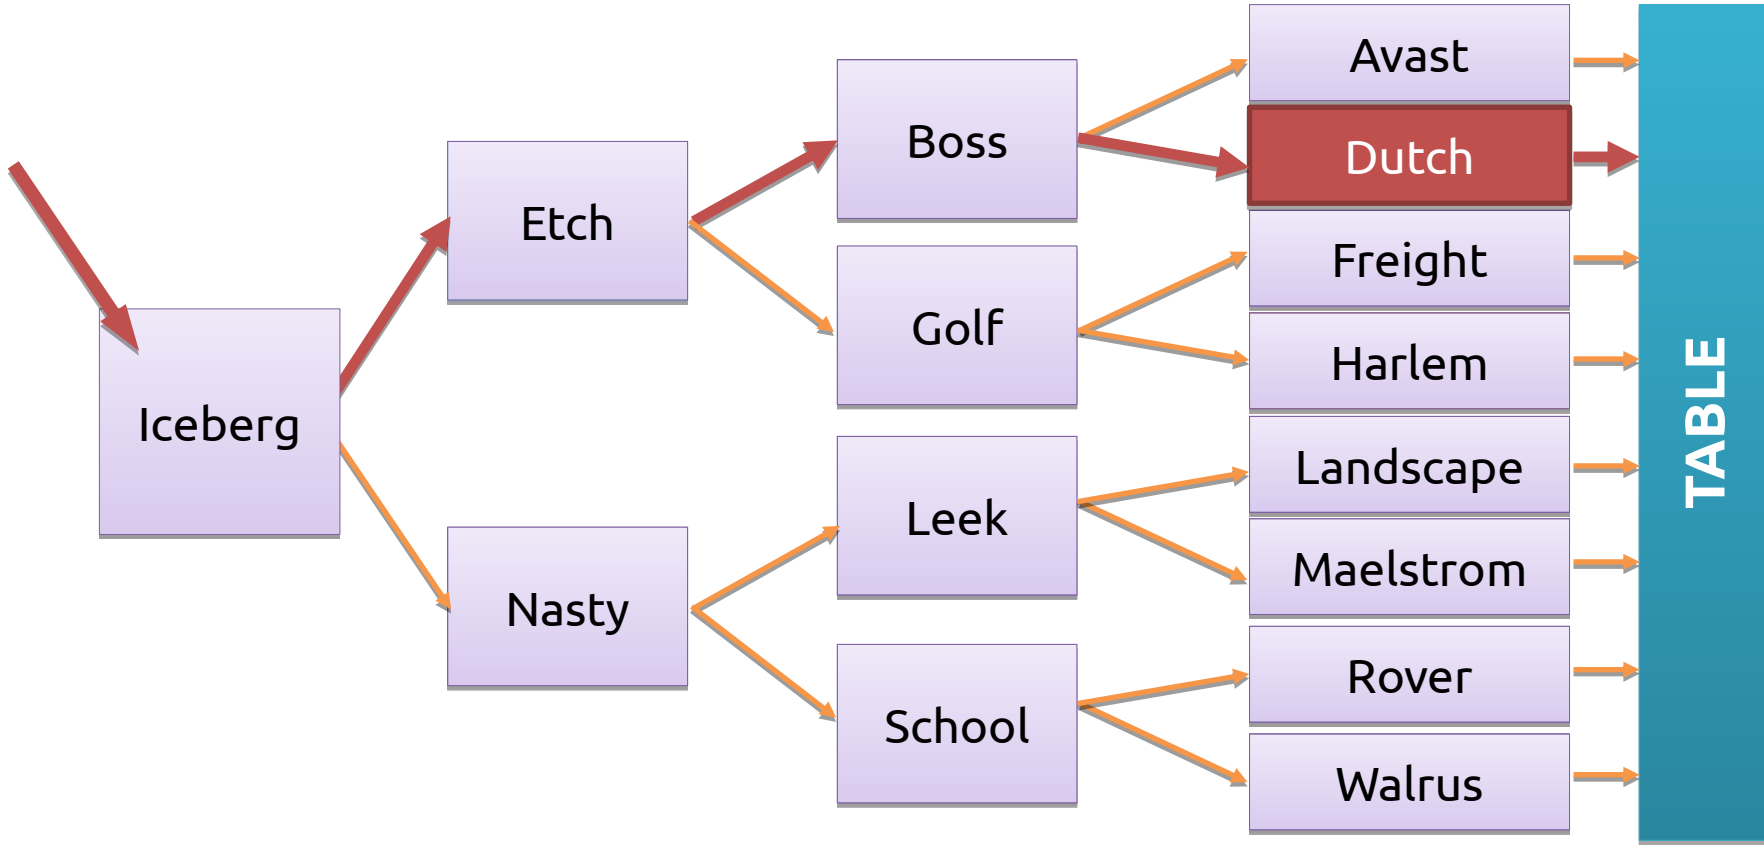
Types of indexes

- BTREE
 - B-TREE in MyISAM, B+TREE in InnoDB
- HASH
 - Only available for MEMORY and NDB
- FULLTEXT
 - Inverted indexes in MyISAM and InnoDB
- SPATIAL
 - RTREEs in MyISAM and InnoDB

Finding “Dutch” with a BTREE



Finding “Dutch” with a BTREE



Do indexes always work? (1/2)

- Can we use an index to make this query faster?

```
SELECT * FROM page WHERE page_title like 'Spa%';
```

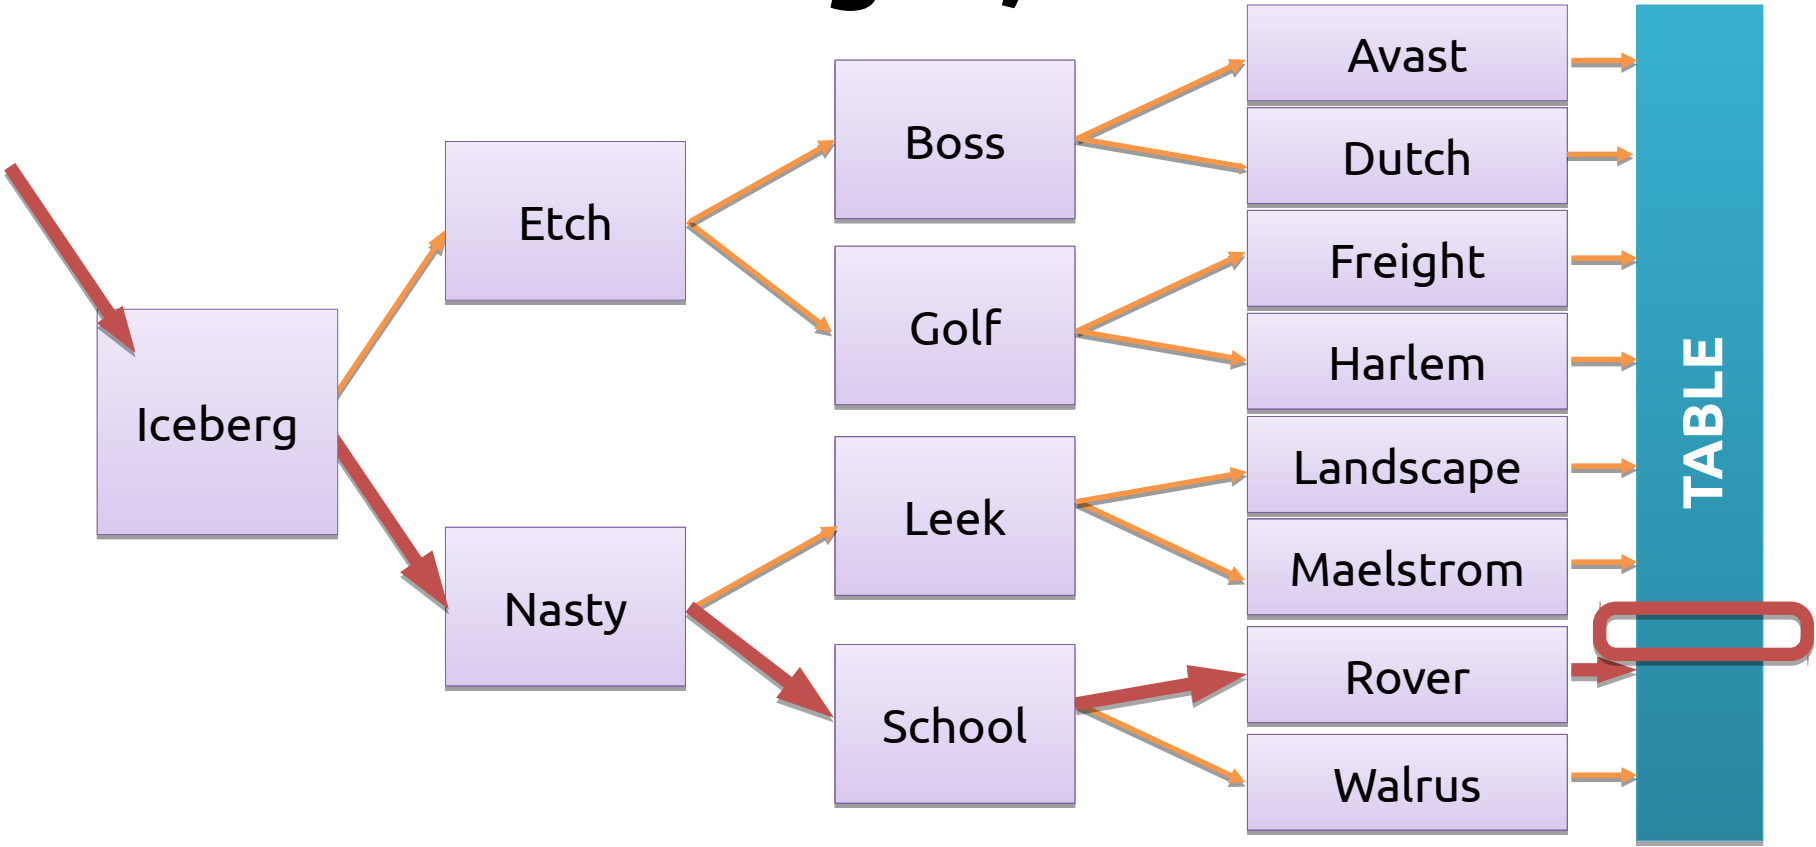
It is a range

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page WHERE  
page_title like 'Spa%'\G
```

```
***** 1. row *****  
      id: 1  
  select_type: SIMPLE  
      table: page  
      type: range  
possible_keys: page_title  
      key: page_title  
   key_len: 257  
      ref: NULL  
     rows: 94  
  Extra: Using index condition  
1 row in set (0.00 sec)
```

Despite not being an equality, we can use the index to find the values quickly

BTREE Indexes can be used for ranges, too



Do indexes always work? (2/2)

- What about this other query?

```
SELECT * FROM page WHERE page_title like '%utch';
```

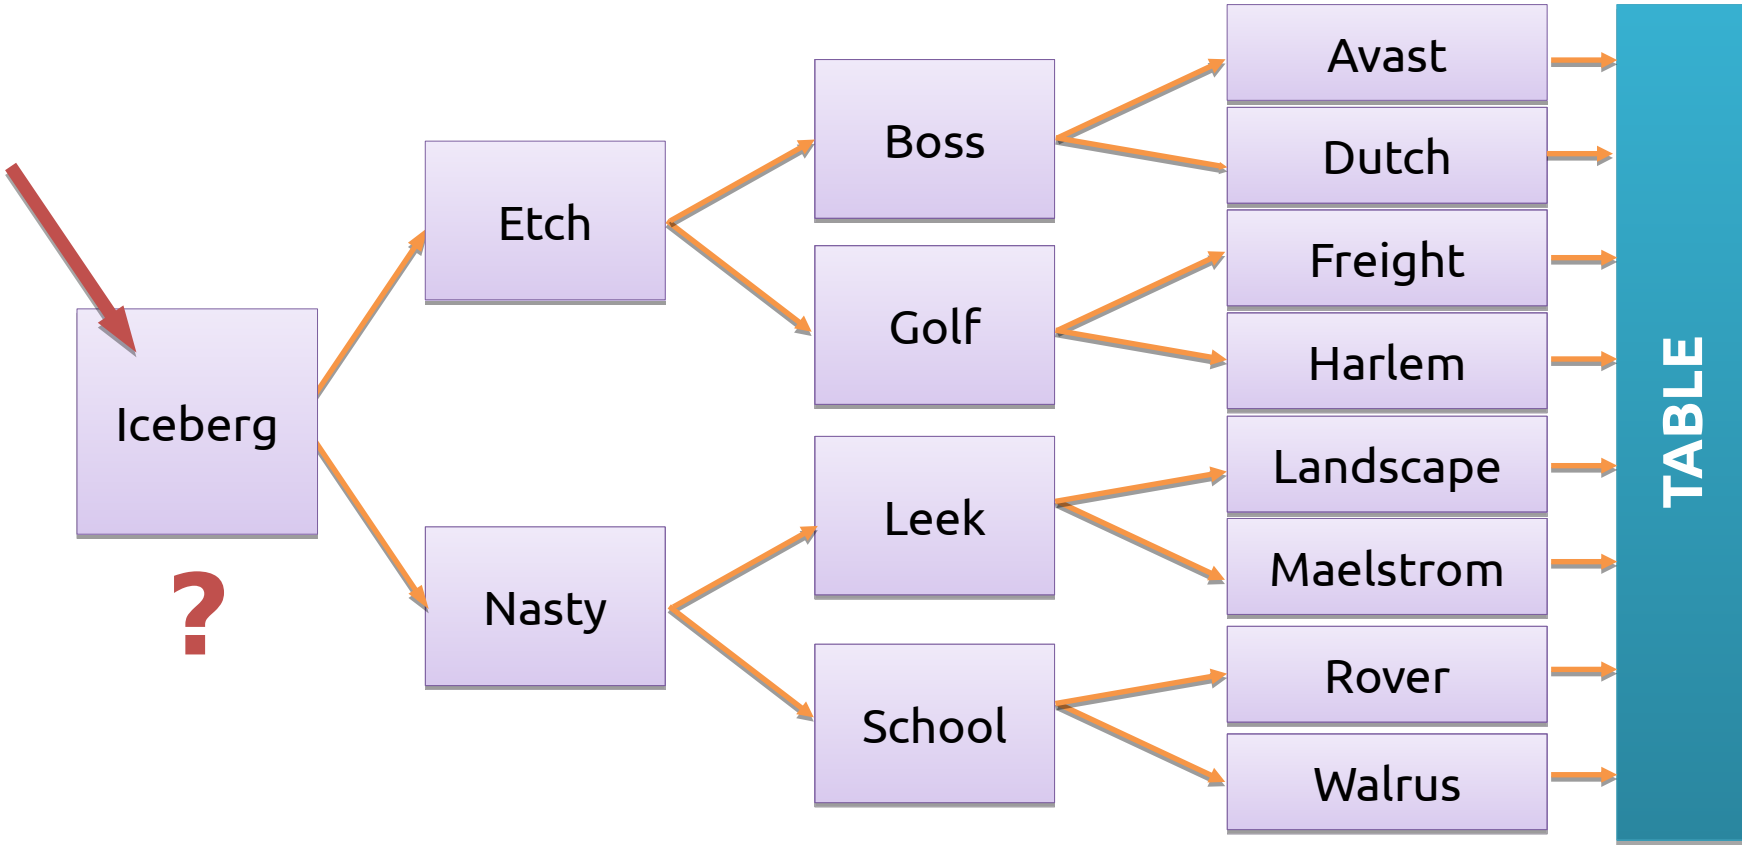
Let's check with EXPLAIN

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page WHERE  
page_title like '%utch'\G
```

```
***** 1. row *****  
      id: 1  
    select_type: SIMPLE  
      table: page  
       type: ALL  
possible_keys: NULL  
       key: NULL  
    key_len: NULL  
       ref: NULL  
      rows: 93189  
    Extra: Using where  
1 row in set (0.00 sec)
```

**No index can be used
for filtering. A full
table scan is
performed.**

BTREE Index



Btree indexes usage

- Filtering
 - Equality (operator '=')
 - Ranges (BETWEEN ... AND, >, <, >=, <=, like 'prefix%')
 - “EXISTS” operators: IN, OR on the same column
- Ordering
 - ORDER BY (indexed columns)
 - GROUP BY (indexed columns)
- Returning values directly from the index
 - Covering index
 - Functions like max(), min(), etc.

type: const

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision WHERE  
rev_id = 2\G
```

```
***** 1. row *****  
      id: 1  
  select_type: SIMPLE  
      table: revision  
      type: const  
possible_keys: PRIMARY  
      key: PRIMARY  
     key_len: 4  
       ref: const  
      rows: 1  
   Extra:  
1 row in set (0.00 sec)
```

'const' is a special case of 'ref', when the index can assure that only 1 results can be returned (equality + primary key or unique key). It is faster.

type: NULL

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision WHERE  
rev_id = -1\G
```

```
***** 1. row *****
```

```
      id: 1  
select_type: SIMPLE  
      table: NULL  
       type: NULL  
possible_keys: NULL  
        key: NULL  
   key_len: NULL  
      ref: NULL  
     rows: NULL
```

**'NULL' is not really a plan,
just an optimization that
allow discarding
immediately impossible
conditions**


```
Extra: Impossible WHERE noticed after reading const
```

```
tables
```

```
1 row in set (0.00 sec)
```

type: ref_or_null


```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM user WHERE  
user_email_token = '0' OR user_email_token IS NULL\G  
***** 1. row *****  
      id: 1  
  select_type: SIMPLE  
      table: user  
      type: ref_or_null  
possible_keys: user_email_token  
      key: user_email_token  
    key_len: 33  
      ref: const  
     rows: 2  
  Extra: Using index condition; Using where  
1 row in set (0.00 sec)
```



**Equivalent to 'ref', but
also has into account
NULL values**

type: range (using IN / OR)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page WHERE
page_title IN ('Dutch', 'English', 'Spanish')\G
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: page
         type: range
possible_keys: page_title
          key: page_title
        key_len: 257
         ref: NULL
         rows: 4
   Extra: Using index condition
1 row in set (0.00 sec)
```



Despite being a range, its execution is very different from ranges using like, between or inequality operators

Is this a bug? (1/2)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page WHERE  
page_namespace = 2\G
```


```
***** 1. row *****  
      id: 1  
    select_type: SIMPLE  
      table: page  
       type: ref  
possible_keys: name_title  
       key: name_title  
    key_len: 4  
       ref: const  
      rows: 45  
    Extra:  
1 row in set (0.00 sec)
```

**An index is used to
return pages with ns=2**

Is this a bug? (2/2)

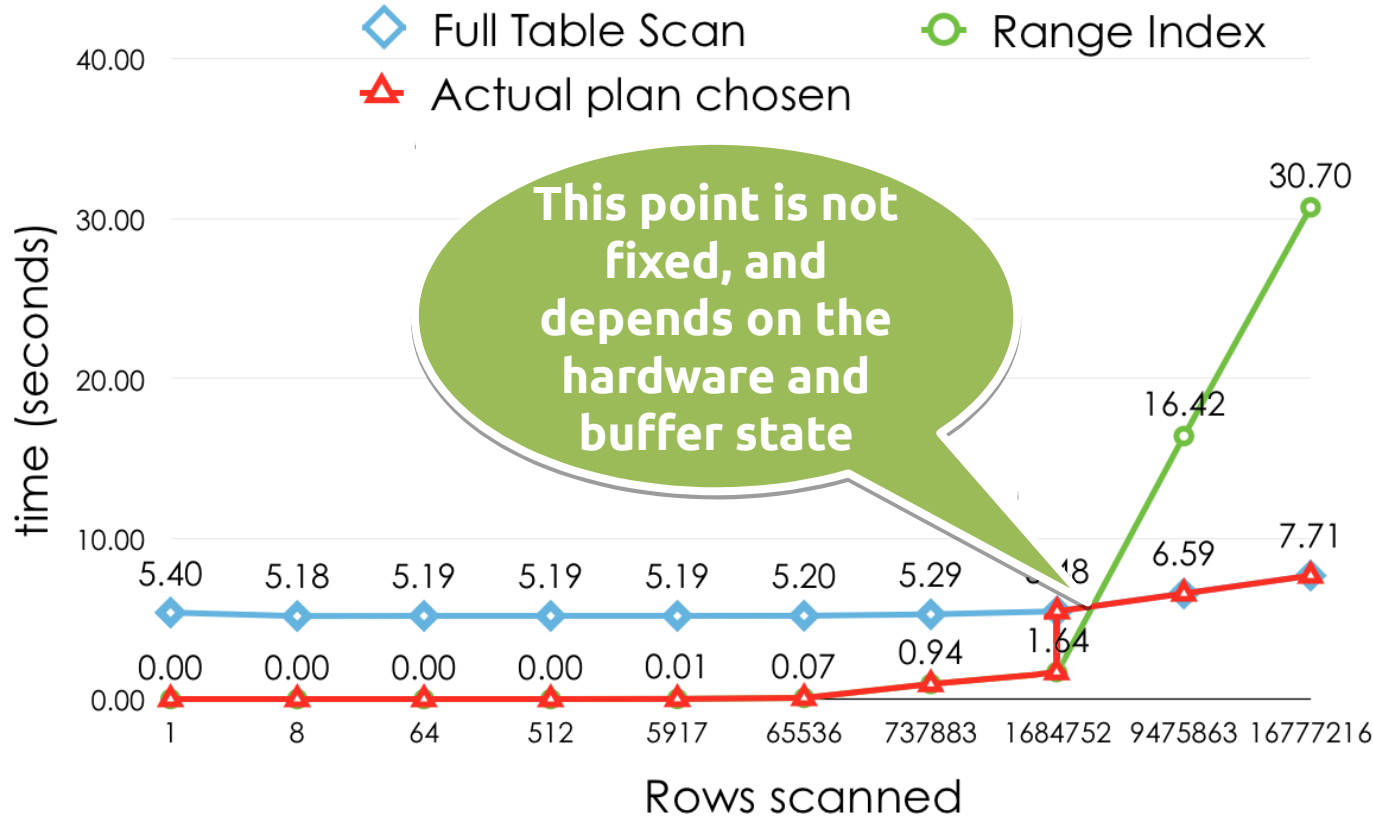
```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page WHERE  
page_namespace = 0\G
```

```
***** 1. row *****  
      id: 1  
    select_type: SIMPLE  
      table: page  
       type: ALL  
possible_keys: name_title  
         key: NULL  
    key_len: NULL  
       ref: NULL  
      rows: 7493  
    Extra: Using where  
1 row in set (0.00 sec)
```



The index is not used
with ns=0

Using an index is sometimes suboptimal



What index should we add to make this query faster?

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision WHERE  
left(rev_timestamp, 6) = '201509'\G
```

```
***** 1. row *****  
      id: 1  
  select_type: SIMPLE  
        table: revision  
         type: ALL  
possible_keys: NULL  
          key: NULL  
       key_len: NULL  
         ref: NULL  
        rows: 163253  
   Extra: Using where  
1 row in set (0.00 sec)
```

The table has already an index on **rev_timestamp**

```
MariaDB [nlwiktionary]> SHOW CREATE TABLE revision\G
```

```
***** 1. row *****
```

```
Table: revision
```

```
Create Table: CREATE TABLE `revision` (  
...
```

```
KEY `rev_timestamp` (`rev_timestamp`),
```



We need to rewrite the query

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision WHERE
rev_timestamp >= '201509' and rev_timestamp < '201510'\G
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: revision
         type: range
possible_keys: rev_timestamp
          key: rev_timestamp
       key_len: 14
         ref: NULL
        rows: 7
   Extra: Using index condition
1 row in set (0.00 sec)
```

This transformation is not trivial or even possible in all cases

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision WHERE
substr(rev_timestamp, 5, 2) = '09'\G
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: revision
         type: ALL
possible_keys: NULL
          key: NULL
        key_len: NULL
         ref: NULL
         Rows: 173154
       Extra: Using where
1 row in set (0.00 sec)
```

Can you think a way to
improve this query?

Indexes for Ordering

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page ORDER BY
page_touched DESC LIMIT 10\G
***** 1. row *****
      id: 1
    select_type: SIMPLE
        table: page
         type: ALL
possible_keys: NULL
         key: NULL
      key_len: NULL
         ref: NULL
        rows: 8720
   Extra: Using filesort
1 row in set (0.00 sec)
```

"Using filesort" indicates that an ordering is needed before returning the results

If that is frequent, we can create an index on page_touched...

```
[nlwiktionary]> ALTER TABLE page ADD INDEX page_page_touched(page_touched);
Query OK, 0 rows affected (0.30 sec)
Records: 0 Duplicates: 0 Warnings: 0
```

```
[nlwiktionary]> EXPLAIN SELECT * FROM page ORDER BY page_touched DESC LIMIT 10\G
```

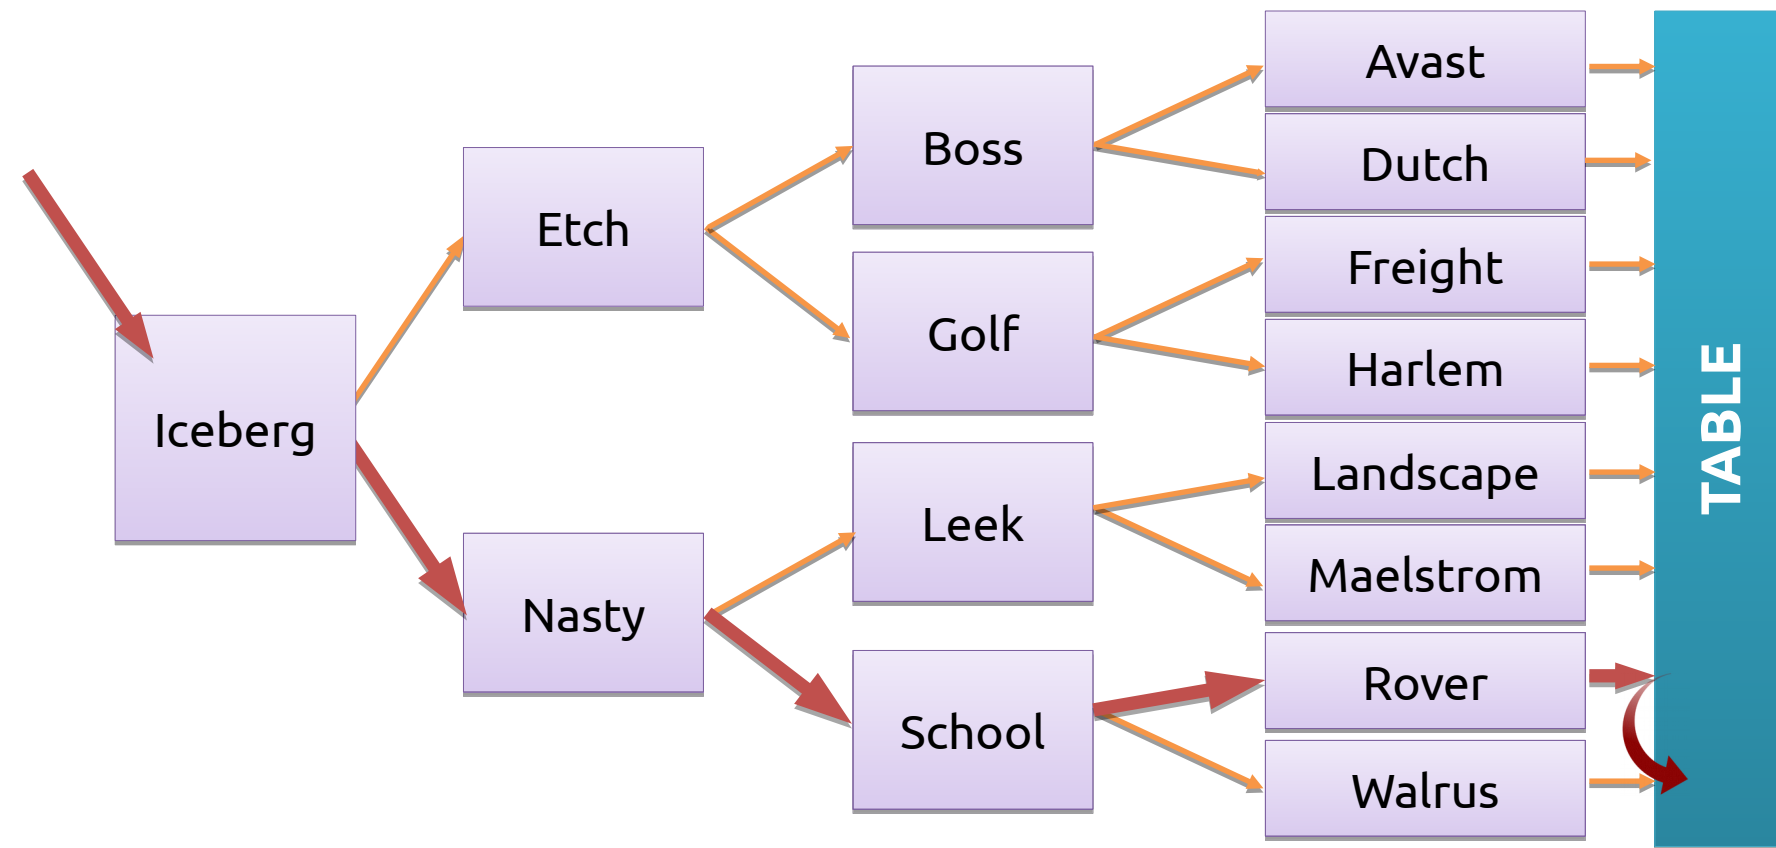
```
***** 1. row *****
```

```
      id: 1
  select_type: SIMPLE
        table: page
         type: index
possible_keys: NULL
          key: page_page_touched
      key_len: 14
         ref: NULL
        rows: 10
      Extra:
1 row in set (0.00 sec)
```

The index does not produce any advantage for filtering

However, it is very effective by helping avoiding the sort phase

It can return data in index order faster



Indexes and GROUP BY (no indexes)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT rev_page, count(*) FROM
revision IGNORE INDEX(rev_page_id, page_timestamp,
page_user_timestamp) GROUP BY rev_page\G
```

```
***** 1. row *****
```

```
id: 1
```

```
select_type: SIMPLE
```

```
table: revision
```

```
type: ALL
```

```
possible_keys: NULL
```

```
key: NULL
```

```
key_len: NULL
```

```
ref: NULL
```

```
rows: 201094
```

```
Extra: Using temporary; Using filesort
```

```
1 row in set (0.00 sec)
```

Without indexes, a temporary table is created to order results

Trick: ORDER BY NULL avoids filesort

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision GROUP BY substr(rev_timestamp, 5, 2) = '09'\G
```

```
***** 1. row *****
```

```
id: 1
select_type: SIMPLE
table: revision
type: ALL
possible_keys: NULL
key: NULL
key_len: NULL
ref: NULL
rows: 196824
Extra: Using temporary; Using filesort
1 row in set (0.00 sec)
```

There is no good index in this case

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision GROUP BY substr(rev_timestamp, 5, 2) = '09' ORDER BY NULL\G
```

```
***** 1. row *****
```

```
...
rows: 196871
Extra: Using temporary
1 row in set (0.00 sec)
```

The advantage is not too big, but it avoids the filesort

Indexes and GROUP BY (rev_page_id)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT rev_page, count(*) FROM  
revision GROUP BY rev_page\G
```

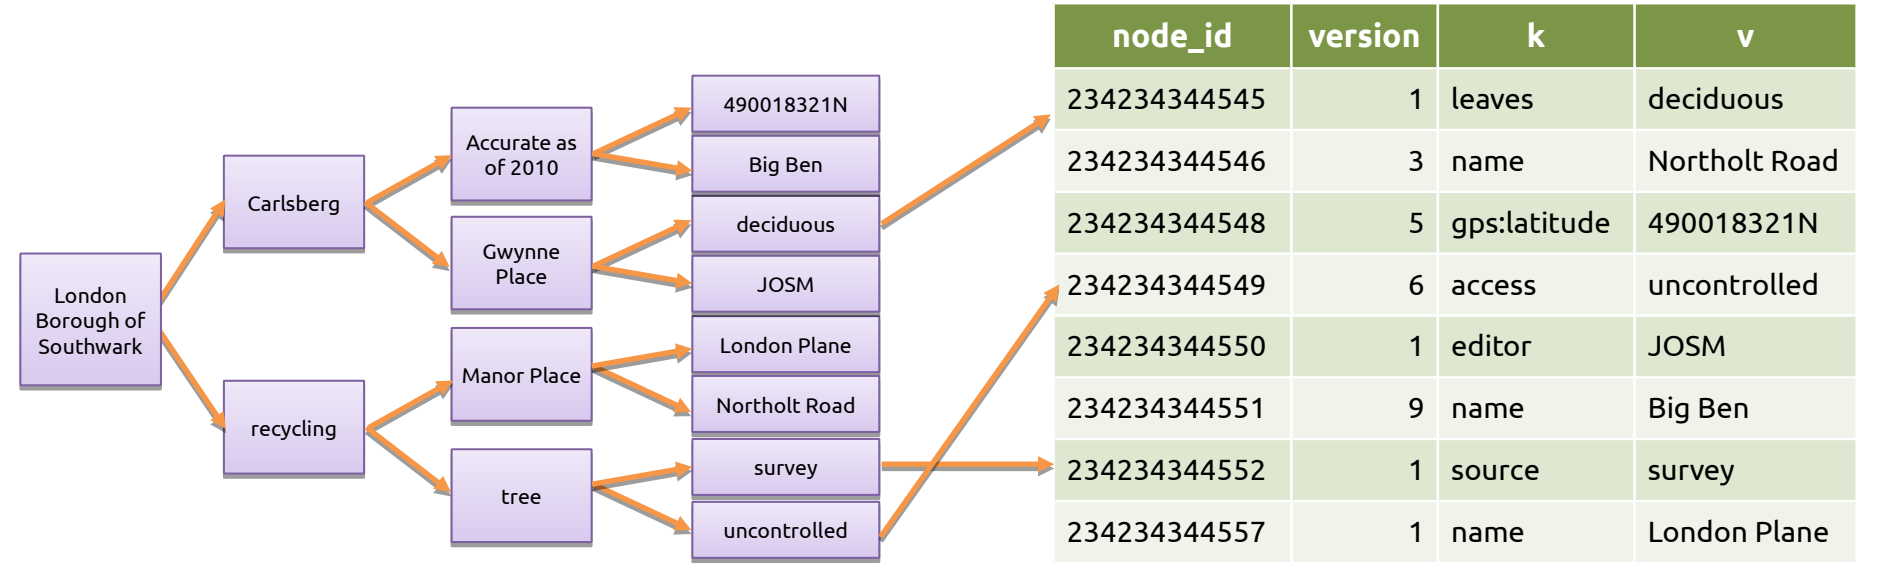
```
***** 1. row *****
```

```
id: 1  
select_type: SIMPLE  
table: revision  
type: index  
possible_keys: NULL  
key: rev_page_id  
key_len: 8  
ref: NULL  
rows: 192388  
Extra: Using index  
1 row in set (0.00 sec)
```

The index does not produce any advantage for filtering (there is no WHERE clause)

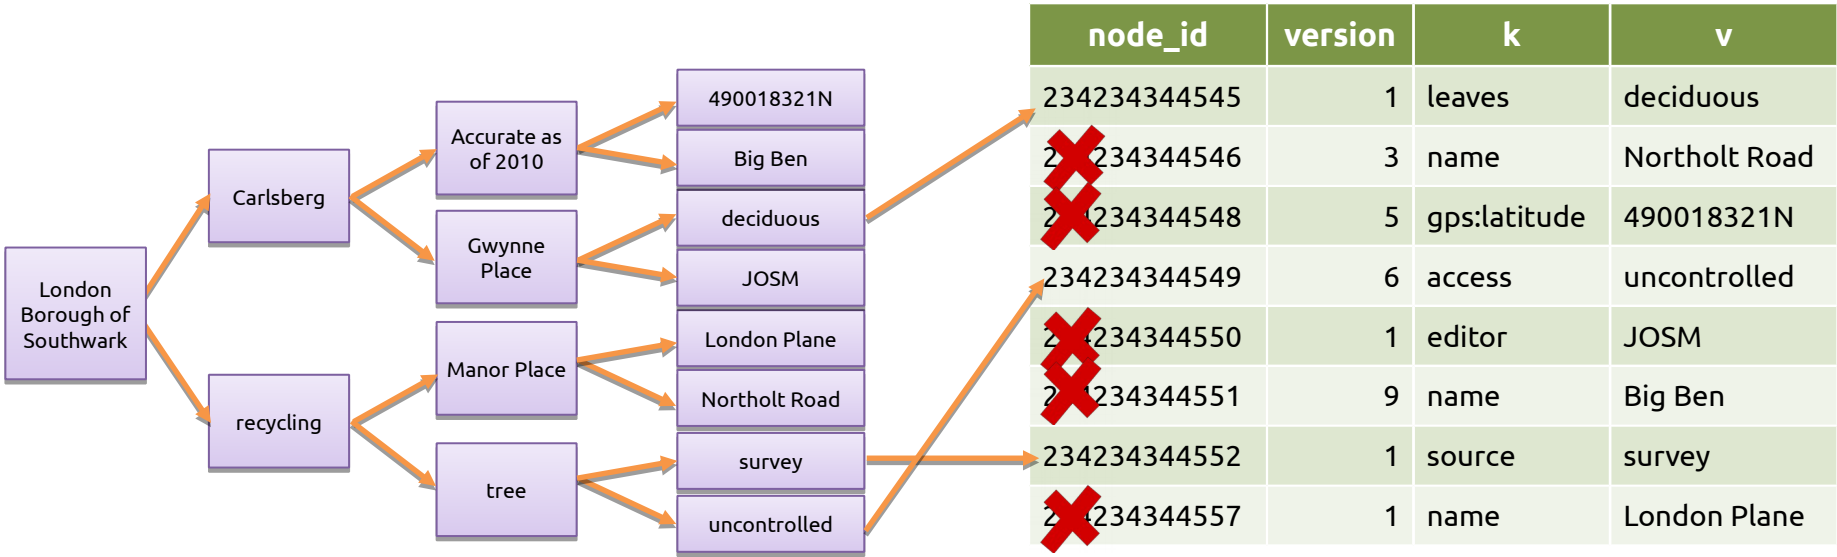
However, thanks to it we avoid a sort and a temporary table

type: index, loose index scan and covering index (1/3)



With 'type:index', all rows are read in index order (full index scan)

type: index, loose index scan and covering index (2/3)

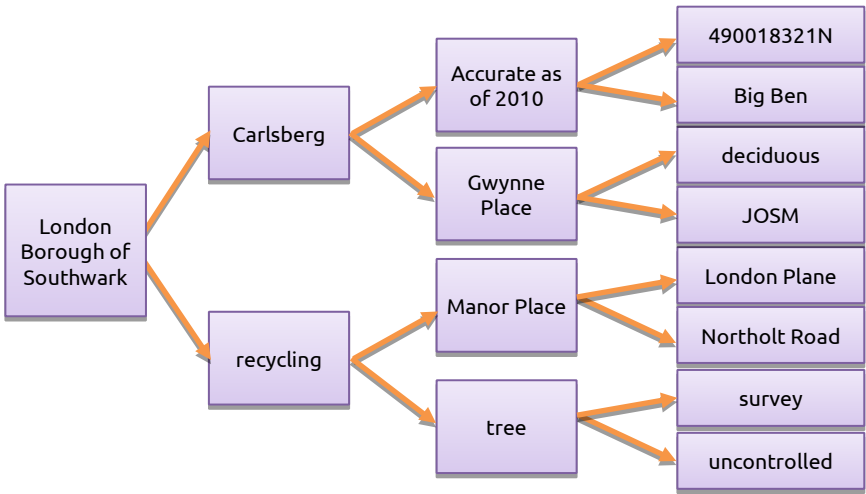


If we have in addition 'Using index for group-by' we have the loose index scan optimization

Loose Index Scan Example

```
MariaDB [nlwiktionary]> EXPLAIN SELECT rev_page,  
max(rev_timestamp) FROM revision GROUP BY rev_page\G  
***** 1. row *****  
      id: 1  
    select_type: SIMPLE  
      table: revision  
        type: range  
possible_keys: NULL  
         key: page_timestamp  
    key_len: 4  
         ref: NULL  
       rows: 9769  
  Extra: Using index for group-by  
1 row in set (0.00 sec)
```


type: index, loose index scan and covering index (3/3)



node	version	v
234234344	1	deciduous
234234344		Northolt Road
234234344548		490018321N
234234344549		uncontrolled
234234344		JOSM
234234344	9	Big Ben
234234344	1	survey
234234344	1	London Plane

If we have in addition 'Using index' we have the covering index optimization

Covering Index Example (1/3)

```
MariaDB [nlwiktionary]> ALTER TABLE revision DROP INDEX rev_page_id, drop index page_timestamp,
drop index page_user_timestamp;
Query OK, 0 rows affected (0.05 sec)
Records: 0 Duplicates: 0 Warnings: 0
```

Let's start with no indexes

```
MariaDB [nlwiktionary]> EXPLAIN SELECT count(DISTINCT rev_user) FROM revision WHERE rev_page =
790\G
```

```
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: revision
         type: ALL
possible_keys: NULL
          key: NULL
         key_len: NULL
          ref: NULL
          rows: 218384
     Extra: Using where
1 row in set (0.00 sec)
```

```
MariaDB [nlwiktionary]> SELECT count(DISTINCT rev_user) FROM revision WHERE rev_page = 790\G
***** 1. row *****
count(DISTINCT rev_user): 1
1 row in set (0.06 sec)
```

Covering Index Example (2/3)

```
MariaDB [nlwiktionary]> ALTER TABLE revision ADD INDEX revision_rev_page(rev_page);
Query OK, 0 rows affected (0.44 sec)
Records: 0 Duplicates: 0 Warnings: 0
```

```
MariaDB [nlwiktionary]> EXPLAIN SELECT count(DISTINCT rev_user) FROM revision WHERE rev_page = 790\G
```

```
***** 1. row *****
```

```
id: 1
select_type: SIMPLE
table: revision
type: ref
possible_keys: revision_rev_page
key: revision_rev_page
key_len: 4
ref: const
rows: 4863
Extra:
1 row in set (0.00 sec)
```

**Adding an index on rev_page
increases the speed due to
improved filtering**

```
MariaDB [nlwiktionary]> SELECT count(DISTINCT rev_user) FROM revision WHERE rev_page = 790\G
```

```
***** 1. row *****
```

```
count(DISTINCT rev_user): 1
1 row in set (0.01 sec)
```

Covering Index Example (3/3)

```
MariaDB [nlwiktionary]> ALTER TABLE revision ADD INDEX revision_rev_page_rev_user(rev_page, rev_user);
Query OK, 0 rows affected (1.48 sec)
Records: 0 Duplicates: 0 Warnings: 0
```

```
MariaDB [nlwiktionary]> EXPLAIN SELECT count(DISTINCT rev_user) FROM revision WHERE rev_page = 790\G
```

```
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: revision
         type: ref
possible_keys: revision_rev_page,revision_rev_page_rev_u
            key: revision_rev_page_rev_user
         key_len: 4
          ref: const
         rows: 4863
      Extra: Using index
1 row in set (0.00 sec)
```

rev_page, rev_user does not increase the index selectiveness, but allow to return results directly from the index

```
MariaDB [nlwiktionary]> SELECT count(DISTINCT rev_user) FROM revision WHERE rev_page = 790\G
***** 1. row *****
count(DISTINCT rev_user): 1
1 row in set (0.00 sec)
```

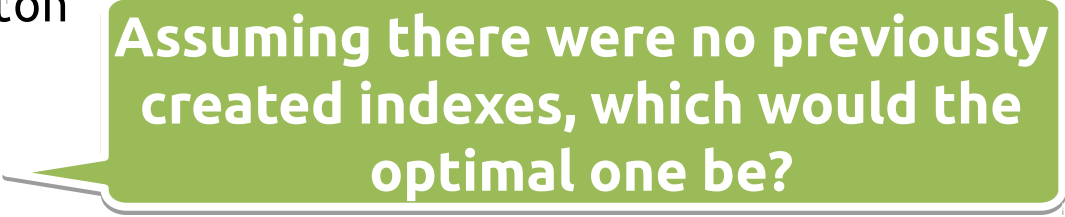
The speed difference can be huge

Query Optimization: From 0 to 10 (and up to 5.7)

MULTI-COLUMN INDEXES

In many cases, conditions are applied on more than one column

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision WHERE  
rev_page = 790 and rev_timestamp < '2008'\G  
***** 1. row *****  
      id: 1  
    select_type: SIMPLE  
      table: revision  
        type: ALL  
possible_keys: NULL  
         key: NULL  
    key_len: NULL  
         ref: NULL  
       rows: 686822  
    Extra: Using where  
1 row in set (0.00 sec)
```



Assuming there were no previously created indexes, which would the optimal one be?

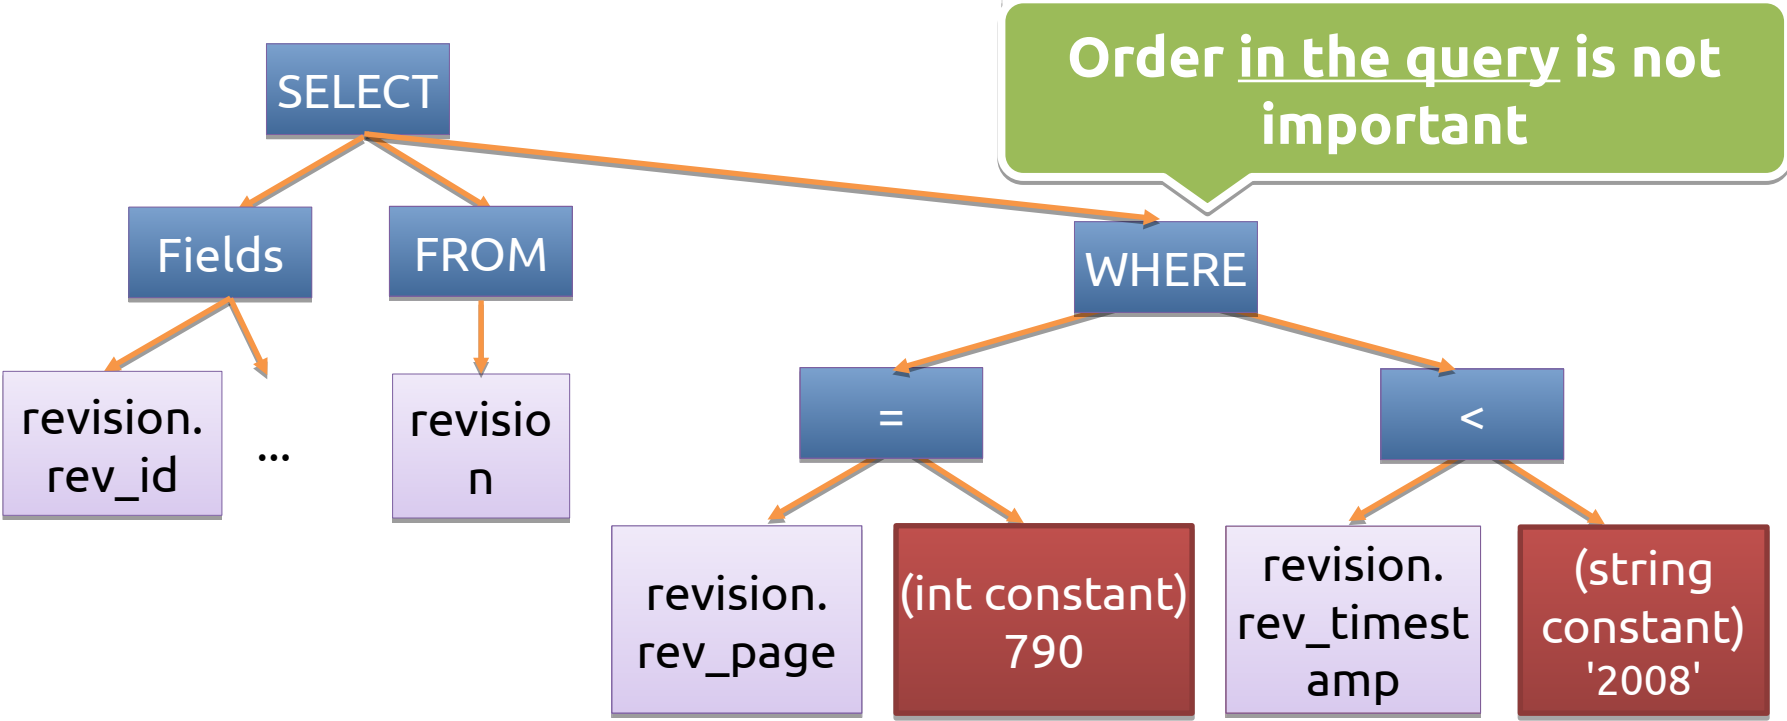
Options for indexes

- 1 index on column (rev_page)
- 1 index on column (rev_timestamp)
- 2 indexes, 1 on (rev_page) and another on (rev_timestamp)
- 1 multi-column index on (rev_page, rev_timestamp)
- 1 multi-column index on (rev_timestamp, rev_page)

**Are these last 2 different from each other?
Would it depend on the query order?**

A brief reminder about query parsing

```
SELECT * FROM revision WHERE rev_page = 790 and rev_timestamp < '2008'
```



Index on (rev_page)

```
MariaDB [nlwiktionary]> ALTER TABLE revision ADD INDEX rev_page (rev_page);  
Query OK, 0 rows affected (2.31 sec)  
Records: 0 Duplicates: 0 Warnings: 0
```

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision WHERE rev_page = 790 and  
rev_timestamp < '2008'\G
```

```
***** 1. row *****
```

```
id: 1  
select_type: SIMPLE  
table: revision  
type: ref  
possible_keys: rev_page  
key: rev_page  
key_len: 4  
ref: const  
rows: 4863  
Extra: Using where  
1 row in set (0.00 sec)
```

Query time improves
significantly with this
index

Less rows are
scanned

Adding (rev_timestamp)

```
MariaDB [nlwiktionary]> ALTER TABLE revision ADD INDEX rev_timestamp  
(rev_timestamp);  
Query OK, 0 rows affected (1.77 sec)  
Records: 0 Duplicates: 0 Warnings: 0
```

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision WHERE rev_page = 790  
and rev_timestamp < '2008'\G  
***** 1. row *****
```

```
      id: 1  
select_type: SIMPLE  
      table: revision  
      type: ref  
possible_keys: rev_page,rev_timestamp  
      key: rev_page  
      key_len: 4  
      ref: const  
      rows: 4863  
Extra: Using where  
1 row in set (0.01 sec)
```

**In general, only one index can
be used per table access**

**rev_page is preferred
over rev_timestamp**

Forcing the use of (rev_timestamp)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision FORCE  
INDEX(rev_timestamp) WHERE rev_page = 790 and rev_timestamp <  
'2008'\G
```

Forcing the index is worse than type:ALL!

```
id: 1  
select_type: SIMPLE  
table: revision  
type: range  
possible_keys: rev_timestamp  
key: rev_timestamp  
key_len: 14  
ref: NULL  
rows: 343411
```

It is a range access

A lot more accessed rows


```
Extra: Using index condition; Using where  
1 row in set (0.00 sec)
```

Adding (rev_page, rev_timestamp)

```
MariaDB [nlwiktionary]> ALTER TABLE revision ADD INDEX  
rev_page_rev_timestamp(rev_page, rev_timestamp);  
Query OK, 0 rows affected (1.59 sec)  
Records: 0 Duplicates: 0 Warnings: 0
```

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision WHERE rev_page = 790  
and rev_timestamp < '2008'\G
```

```
***** 1. row *****  
      id: 1  
  select_type: SIMPLE  
      table: revision  
       type: range  
possible_keys: rev_page,rev_timestamp,rev_page_rev_timestamp  
          key: rev_page_rev_timestamp  
      key_len: 18  
         ref: NULL  
       rows: 1048  
  Extra: Using index condition  
1 row in set (0.00 sec)
```



**Reduced number of
rows scanned**

Is (rev_timestamp, rev_page) a better option?

```
MariaDB [nlwiktionary]> ALTER TABLE revision ADD INDEX rev_timestamp_rev_page
(rev_timestamp, rev_page);
Query OK, 0 rows affected (1.76 sec)
Records: 0  Duplicates: 0  Warnings: 0
```

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision WHERE rev_page = 790 and
rev_timestamp < '2008'\G
```


```
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: revision
         type: range
possible_keys: rev_page,rev_timestamp,rev_page_rev_timestamp,rev_timestamp_rev_page
           key: rev_page_rev_timestamp
        key_len: 18
           ref: NULL
          rows: 1048
     Extra: Using index condition
1 row in set (0.00 sec)
```

Previous index is still preferred, why?

Forcing (rev_timestamp, rev_page)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision FORCE  
INDEX(rev_timestamp_rev_page) WHERE rev_page = 790 and  
rev_timestamp < '2008'\G
```

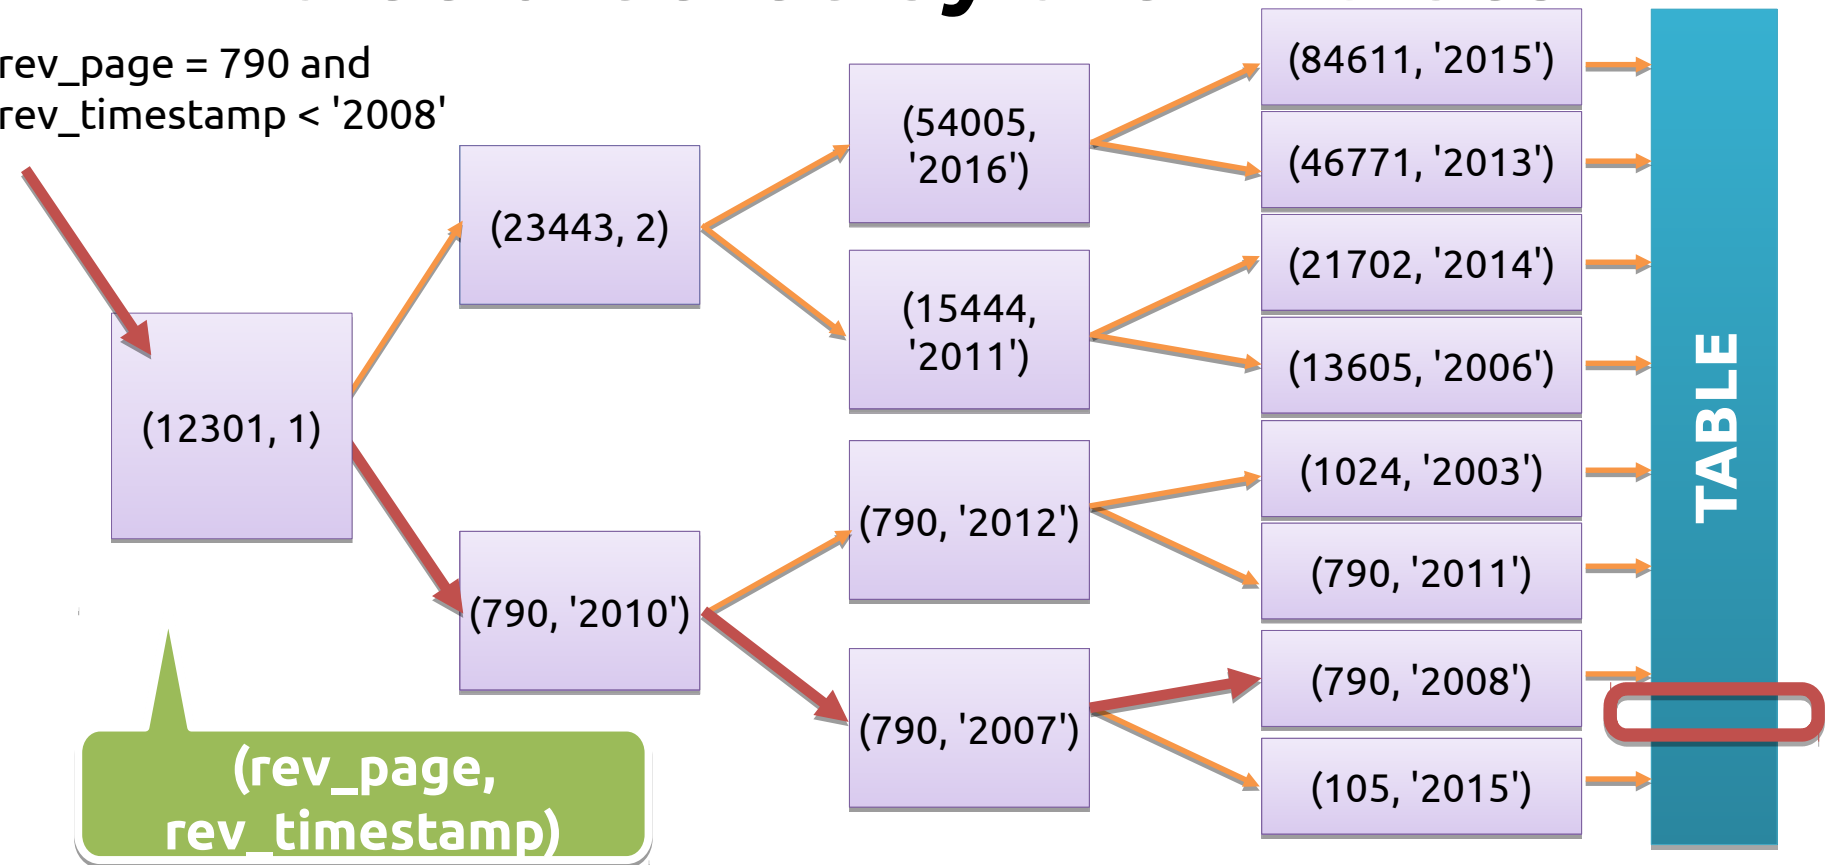
```
***** 1. row *****  
      id: 1  
  select_type: SIMPLE  
        table: revision  
         type: range  
possible_keys: rev_timestamp_rev_page  
          key: rev_timestamp_rev_page  
     key_len: 18  
         ref: NULL  
        rows: 343411  
   Extra: Using index condition  
1 row in set (0.00 sec)
```



Only the first column is being
used effectively for filtering

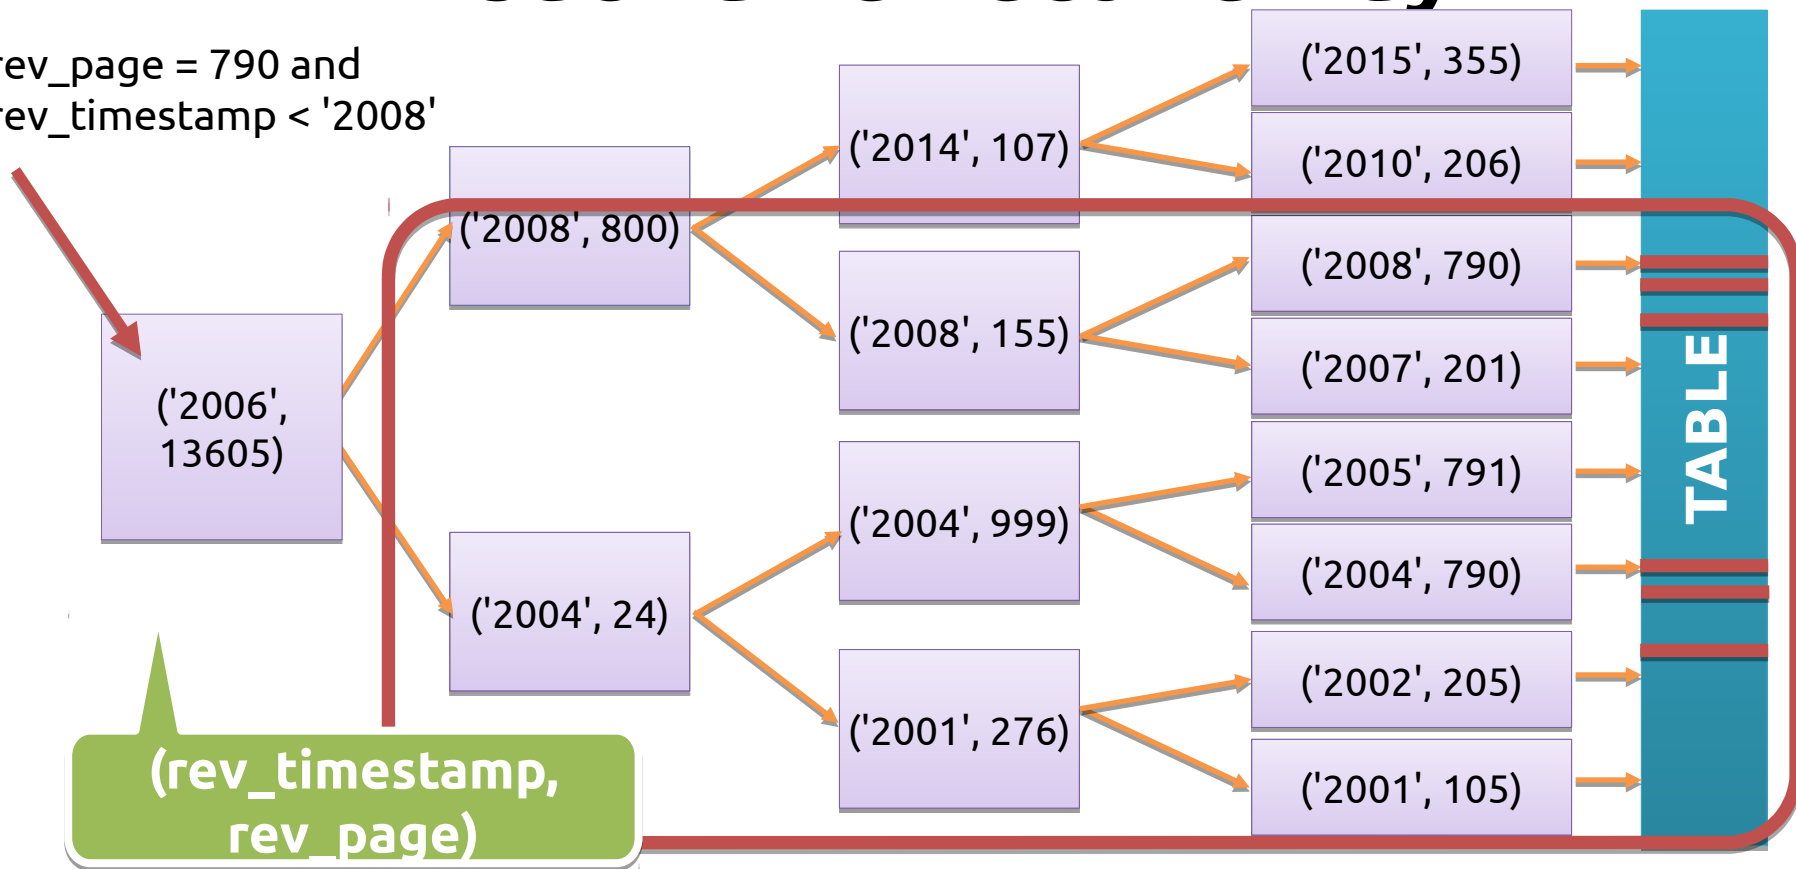
A compound index produces a single tree ordered by the 2 values

rev_page = 790 and
rev_timestamp < '2008'



The alternative index cannot be used in such an effective way

rev_page = 790 and
rev_timestamp < '2008'



Order and column selection

- Range access using $>$, $<$, $>=$, $<=$, BETWEEN can only be filtered once effectively, at the end of an index
- When selecting indexes, prefer columns with high cardinality (very selective)
 - The optimal index can depend on the constants used

Use of "Handler_*" statistics

- They are post-execution statistics at row level
 - Unlike EXPLAIN's "rows" column, they are exact, not a guess
 - They allow to compare query execution performance in a deterministic way, independently of the execution time

"Handler" Stats (indexed)

```
mysql> SHOW SESSION STATUS like 'Hand%';
```

Variable_name	Value
Handler_commit	80
Handler_delete	0
Handler_discover	0
Handler_external_lock	166
Handler_mrr_init	0
Handler_prepare	0
Handler_read_first	23
Handler_read_key	736212
Handler_read_last	0
Handler_read_next	22208001
Handler_read_prev	0
Handler_read_rnd	665215
Handler_read_rnd_next	14223297
Handler_rollback	0
Handler_savepoint	0
Handler_savepoint_rollback	0
Handler_update	66970
Handler_write	2869409

Number of times that the first entry of an index was read. It may indicate the number of full index scans

Number of time a row has been retrieved using an index

Next row has been requested in index order (typical for index scans or ranges)

18 rows in set (0.00 sec)

"Handler" Stats (unindexed)

```
mysql> SHOW SESSION STATUS like 'Hand%';
```

Variable_name	Value
Handler_commit	80
Handler_delete	0
Handler_discover	0
Handler_external_lock	166
Handler_mrr_init	0
Handler_prepare	0
Handler_read_first	23
Handler_read_key	736212
Handler_read_last	0
Handler_read_next	22208001
Handler_read_prev	0
Handler_read_rnd	665215
Handler_read_rnd_next	14223297
Handler_rollback	0
Handler_savepoint	0
Handler_savepoint_rollback	0
Handler_update	66970
Handler_write	2869409

18 rows in set (0.00 sec)

A row has been requested in a specific position (typical for joins or order by without indexes)

Request to read the next row in "table order" (typical for full table scans)

Insertions in SELECTS may indicate temporary tables

Comparing statistics of the previous indexes (no indexes)

```
MariaDB [nlwiktionary]> FLUSH STATUS;  
Query OK, 0 rows affected (0.00 sec)
```

```
MariaDB [nlwiktionary]> SELECT * FROM revision IGNORE INDEX(rev_page, rev_timestamp, rev_page_rev_timestamp,  
rev_timestamp_rev_page) WHERE rev_page = 790 and rev_timestamp < '2008';  
1049 rows in set (0.58 sec)
```

```
MariaDB [nlwiktionary]> SHOW STATUS like 'Hand%';
```

Variable_name	Value
Handler_commit	1
Handler_delete	0
...	
Handler_read_first	0
Handler_read_key	0
Handler_read_last	0
Handler_read_next	0
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_deleted	0
Handler_read_rnd_next	820309
Handler_rollback	0
...	
Handler_update	0
Handler_write	0

25 rows in set (0.00 sec)

Typical result for a
full table scan

Index on (rev_page)

MariaDB [nlwiktionary]> SHOW STATUS like 'Hand%';

Variable_name	Value
Handler_commit	1
Handler_delete	0
...	
Handler_read_first	0
Handler_read_key	1
Handler_read_last	0
Handler_read_next	4864
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_deleted	0
Handler_read_rnd_next	0
Handler_rollback	0
...	
Handler_update	0
Handler_write	0

25 rows in set (0.01 sec)

Using the index, request the first row with rev_page=790

Then, scan them one by one in index order

Index on (rev_timestamp)

MariaDB [nlwiktionary]> SHOW STATUS like 'Hand%';

Variable_name	Value
Handler_commit	1
Handler_delete	0
...	
Handler_read_first	0
Handler_read_key	1
Handler_read_last	0
Handler_read_next	199155
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_deleted	0
Handler_read_rnd_next	0
Handler_rollback	0
...	
Handler_update	0
Handler_write	0

25 rows in set (0.00 sec)

Let's ignore ICP for now

Using the index, request the first row where rev_timestamp<2008

Then, scan them one by one in index order (more are matched)

Index on (rev_page, rev_timestamp)

MariaDB [nlwiktionary]> SHOW STATUS like 'Hand%';

Variable_name	Value
Handler_commit	1
Handler_delete	0
...	
Handler_read_first	0
Handler_read_key	1
Handler_read_last	0
Handler_read_next	1049
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_deleted	0
Handler_read_rnd_next	0
Handler_rollback	0
...	
Handler_update	0
Handler_write	0

25 rows in set (0.00 sec)

With both conditions covered, we can find the actual first row that matches the condition using the index

Rows scanned == Rows returned

Index on (rev_timestamp, rev_page), no ICP

MariaDB [nlwiktionary]> SHOW STATUS like 'Hand%';

Variable_name	Value
Handler_commit	1
Handler_delete	0
...	
Handler_read_first	0
Handler_read_key	1
Handler_read_last	0
Handler_read_next	199155
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_deleted	0
Handler_read_rnd_next	0
Handler_rollback	0
...	
Handler_update	0
Handler_write	0

25 rows in set (0.00 sec)


Assuming no ICP, exact same results than with (rev_timestamp). The extra column does not help. Also, EXPLAIN's row count was very off.

Redundant Indexes

- Creating all 4 previous indexes in production is not a great idea
 - "Left-most index prefix" allows, for example (rev_page, rev_timestamp) doing everything you can do with (rev_page)
 - If two indexes have equal selectivity, MySQL chooses the shortest one

"Left-most index" Example

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision FORCE
INDEX(rev_page_rev_timestamp) WHERE rev_page = 790\G
***** 1. row *****
      id: 1
    select_type: SIMPLE
        table: revision
         type: ref
possible_keys: rev_page_rev_timestamp
          key: rev_page_rev_time
        key_len: 4
           ref: const
          rows: 4863
         Extra:
1 row in set (0.00 sec)
```



Only the first column is used

Duplicate Indexes

- It is very easy to create indexes with the same exact definition (same columns and ordering)
 - Set a convention for index naming (e.g `tablename_column1_column2_idx`) – MySQL does not allow 2 indexes with the same identifier
 - Since MySQL 5.6, an warning is thrown if a duplicate index is created

pt-duplicate-index-checker

```
$ pt-duplicate-key-checker h=127.0.0.1,P=5621,u=msandbox,p=msandbox
[...]
# rev_timestamp is a left-prefix of rev_timestamp_rev_page
# Key definitions:
#   KEY `rev_timestamp` (`rev_timestamp`),
#   KEY `rev_timestamp_rev_page` (`rev_timestamp`,`rev_page`)
# Column types:
#   `rev_timestamp` binary(14) not null default '\0\0\0\0\0\0\0\0\0\0\0\0\0\0\0\0'
#   `rev_page` int(10) unsigned not null
# To remove this duplicate index, execute:
ALTER TABLE `nlwiktionary`.`revision` DROP INDEX `rev_timestamp`;

# rev_page is a left-prefix of rev_page_rev_timestamp
# Key definitions:
#   KEY `rev_page` (`rev_page`),
#   KEY `rev_page_rev_timestamp` (`rev_page`,`rev_timestamp`),
# Column types:
#   `rev_page` int(10) unsigned not null
#   `rev_timestamp` binary(14) not null default '\0\0\0\0\0\0\0\0\0\0\0\0\0\0\0\0'
# To remove this duplicate index, execute:
ALTER TABLE `nlwiktionary`.`revision` DROP INDEX `rev_page`;

# ##### Summary of indexes
#####
# Size Duplicate Indexes    15478317
# Total Duplicate Indexes   4
# Total Indexes             285
```

Simple tool to check
redundant and
duplicate indexes

"OR"-style conditions over the same column

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision WHERE  
rev_page = 790 OR rev_page = 795 OR rev_page = 1024\G  
***** 1 row *****  
      id: 1  
    select_type: SIMPLE  
        table: revision  
         type: range  
possible_keys: rev_page,rev_page_rev_timestamp  
          key: rev_page  
       key_len: 4  
         ref: NULL  
        rows: 4890  
   Extra: Using index condition; Using where  
1 row in set (0.01 sec)
```

Equivalent to:

**SELECT * FROM revision WHERE
rev_page IN (790, 795, 1024)**

Equivalent to:
SELECT * FROM revision WHERE
rev_page IN (790, 795, 1024)

Handlers on "IN" / "OR" conditions over the same column

Variable_name	Value
Handler_commit	1
Handler_delete	0
...	
Handler_prepare	0
Handler_read_first	0
Handler_read_key	3
Handler_read_last	0
Handler_read_next	4891
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_deleted	0
Handler_read_rnd_next	0
Handler_rollback	0
...	
Handler_update	0
Handler_write	0

25 rows in set (0.00 sec)

Despite identifying themselves as “range”s, the execution is slightly different, one index dive (similar to a ref) is done per value. This can be an issue in conditions with thousands of items.

"OR"-style conditions over different columns

- We cannot use a single index efficiently for both conditions
 - We can scan both conditions separately and mix the results, discarding duplicates
 - Or use an index for one condition and not for the other
 - Index merge allows to use two indexes for a single table access simultaneously

Index Merge Example

```

MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision WHERE
rev_page = 790 or rev_timestamp < '2004'\G
***** 1. row *****
      id: 1
    select_type: SIMPLE
      table: revision
        type: index_merge
possible_keys: rev_page,rev_timestamp,rev_page_rev_timestamp,
rev_timestamp_rev_page
          key: rev_page,rev_timestamp
        key_len: 4,14
         ref: NULL
        rows: 4871
      Extra: Using sort_union(rev_page,rev_timestamp); Using
where
1 row in set (0.00 sec)
    
```

Both indexes are used, then combined using the "union" operation

INDEX_MERGE Issues

- Sometimes it is faster to execute the sentence using UNION:
 - This is specially true with (UNION ALL) since MySQL 5.7, if you do not care about duplicates
- There are also interseccion merges, but multi-column indexes are preferred

Disabling optimizer features (I)

- The optimizer_switch variable allows enabling and disabling globally or per session many query optimizer features:

```
MariaDB [nlwiktionary]> SHOW VARIABLES like 'optimizer_switch'\G
```

```
***** 1. row *****
```

```
Variable_name: optimizer_switch
```

```
Value:
```

```
index_merge=on,index_merge_union=on,index_merge_sort_union=on,index_merge_intersection=on,index_merge_sort_intersection=off,engine_condition_pushdown=off,index_condition_pushdown=on,derived_merge=on,derived_with_keys=on,firstmatch=on,loosescan=on,materiazation=on,in_to_exists=on,semijoin=on,partial_match_rowid_merge=on,partial_match_table_scan=on,subquery_cache=on,mrr=off,mrr_cost_based=off,mrr_sort_keys=off,outer_join_with_cache=on,semijoin_with_cache=on,join_cache_incremental=on,join_cache_hashed=on,join_cache_bka=on,optimize_join_buffer_size=off,table_elimination=on,extended_keys=on,exists_to_in=on
```

```
1 row in set (0.00 sec)
```

Deshabilitar características del optimizador (II)

```
MariaDB [nlwiktionary]> SET optimizer_switch='index_merge_sort_union=off';  
Query OK, 0 rows affected (0.00 sec)
```

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision  
or rev_timestamp < '2004'\G
```

```
***** 1. row *****
```

```
id: 1
```

```
select_type: SIMPLE
```

```
table: revision
```

```
type: ALL
```

```
possible_keys:
```

```
rev_page,rev_timestamp,rev_page_rev_timestamp,rev_timestamp_rev_page
```

```
key: NULL
```

```
key_len: NULL
```

```
ref: NULL
```

```
rows: 686822
```

```
Extra: Using where
```

```
1 row in set (0.00 sec)
```

This will only have effect for the current session.

What happens if we have two ranges?

- As seen previously, we cannot use efficiently two range types on the same table access. Alternatives:
 - Use only one index for the most selective column
 - Use index condition pushdown to get an advantage
 - Change one of the two ranges into a discrete "IN" comparison/bucketize with a new column
 - Use quadtrees or R-TREES (spatial indexing)

Example of Bucketizing (I)

```

MariaDB [nlwiktionary]> EXPLAIN SELECT count(*) FROM revision
WHERE rev_timestamp < '2008' AND rev_len > 5500\G
***** 1. row *****
      id: 1
    select_type: SIMPLE
        table: revision
         type: ALL
possible_keys: rev_timestamp,rev_timestamp_rev_page
          key: NULL
       key_len: NULL
         ref: NULL
        rows: 686822
      Extra: Using where
1 row in set (0.00 sec)
    
```

Looks like only an index on (rev_timestamp) or (rev_len) would be useful as we have 2 ranges.

Example of Bucketizing (II)

```
MariaDB [nlwiktictionary]> ALTER TABLE revision ADD COLUMN  
rev_len_cat int;  
Query OK, 0 rows affected (38.28 sec)  
Records: 0 Duplicates: 0 Warnings: 0
```

```
MariaDB [nlwiktictionary]> UPDATE revision set rev_len_cat =  
IF(rev_len < 10000, rev_len div 1000, 10);  
Query OK, 820308 rows affected (15.19 sec)  
Rows matched: 820308 Changed: 820308 Warnings: 0
```

```
MariaDB [nlwiktictionary]> ALTER TABLE revision ADD INDEX  
rev_len_cat_rev_timestamp (rev_len_cat, rev_timestamp);  
Query OK, 0 rows affected (2.11 sec)  
Records: 0 Duplicates: 0 Warnings: 0
```

Example of Bucketizing (III)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT count(*) FROM revision WHERE  
rev_timestamp < '2008' AND rev_len > 5500 AND rev_len_cat IN (5, 6,  
7, 8, 9, 10)\G
```

```
***** 1. row *****  
      id: 1  
  select_type: SIMPLE  
        table: revision  
         type: range  
possible_keys:  
rev_timestamp,rev_timestamp_rev_page,rev_len_cat_rev_timestamp  
      key: rev_len_cat_rev_timestamp  
    key_len: 19  
       ref: NULL  
      rows: 4442  
  Extra: Using where  
1 row in set (0.00 sec)
```

**We did some transformations
to both the structure and the
query.**

Example of Index Condition Pushdown

```
MariaDB [nlwiktionary]> ALTER TABLE revision ADD INDEX  
rev_len_rev_timestamp(rev_len, rev_timestamp);  
Query OK, 0 rows affected (1.77 sec)  
Records: 0 Duplicates: 0 Warnings: 0
```

```
MariaDB [nlwiktionary]> SET optimizer_switch='index_condition_pushdown=on'; EXPLAIN  
SELECT * FROM revision WHERE rev_timestamp < '2008' AND rev_len > 5500\G  
Query OK, 0 rows affected (0.00 sec)
```

```
***** 1. row *****
```

```
id: 1  
select_type: SIMPLE  
table: revision  
type: range  
possible_keys: rev_timestamp,rev_timestamp  
key: rev_len_rev_timestamp  
key_len: 5  
ref: NULL  
rows: 38744  
Extra: Using index condition  
1 row in set (0.00 sec)
```

Index condition pushdown (ICP) enables the engines to use extra parts of the index while avoiding costly row movements to and from the SQL layer

ICP Issues

- Differences in execution time is more significative when the extra column condition is very selective (getting 5x the original performance)
- ICP is ignored when using covering Index, making the performance worse

Does LIMIT improve the performance? (I)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page ORDER BY page_touched\G
***** 1. row *****
      type: ALL
possible_keys: NULL
      key: NULL
     key_len: NULL
        ref: NULL
       rows: 90956
  Extra: Using filesort
```

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page ORDER BY page_touched LIMIT 10\G
***** 1. row *****
      type: index
possible_keys: NULL
      key: page_page_touched
     key_len: 14
        ref: NULL
       rows: 10
  Extra:
1 row in set (0.00 sec)
```

In some cases it can be essential to allow effective usage of the indexes

Does LIMIT improve the performance? (II)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision ORDER BY rev_comment\G
***** 1. row *****
[...]
```

type:	ALL
possible_keys:	NULL
key:	NULL
key_len:	NULL
ref:	NULL
rows:	817636
Extra:	Using filesort

1 row in set (0.00 sec)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision ORDER BY rev_comment LIMIT 10\G
***** 1. row *****
[...]
```

table:	revision
type:	ALL
possible_keys:	NULL
key:	NULL
key_len:	NULL
ref:	NULL
rows:	817636
Extra:	Using filesort

1 row in set (0.00 sec)

In other cases, it has no effect on the scanned rows (just on the returned ones)

Does LIMIT improve the performance? (I)

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page ORDER BY page_title LIMIT 100\G
```

```
***** 1. row *****
```

```
    type: index
possible_keys: NULL
    key: page_title
    key_len: 257
    ref: NULL
    rows: 100
  Extra:
1 row in set (0.00 sec)
```

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM page ORDER BY page_title LIMIT 10000, 100\G
```

```
***** 1. row *****
```

```
    type: ALL
possible_keys: NULL
    key: NULL
    key_len: NULL
    ref: NULL
    rows: 90956
  Extra: Using filesort
1 row in set (0.00 sec)
```

**In this case,
performance will vary
depending on the
offset (not ideal)**

Can we filter and sort at the same time using indexes?

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision WHERE
rev_comment='' ORDER BY rev_timestamp ASC\G
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: revision
         type: ALL
possible_keys: NULL
          key: NULL
       key_len: NULL
         ref: NULL
         rows: 817636
  Extra: Using where; Using filesort
1 row in set (0.00 sec)
```

This query is slow because a) the full table scan

b) Required sort after filtering

Adding and index on (rev_comment, rev_timestamp)

```
MariaDB [nlwiktionary]> ALTER TABLE revision ADD INDEX  
rev_comment_rev_timestamp (rev_comment, rev_timestamp);  
Query OK, 0 rows affected (3.19 sec)  
Records: 0 Duplicates: 0 Warnings: 0
```

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision WHERE rev_comment=''  
ORDER BY rev_timestamp ASC\G
```

```
***** 1. row *****
```

```
id: 1  
select_type: SIMPLE  
table: revision  
type: ref  
possible_keys: rev_comment_rev_timestamp  
key: rev_comment_rev_timestamp  
key_len: 769  
ref: const  
rows: 266462  
Extra: Using index condition; Using where  
1 row in set (0.00 sec)
```

**Both type: ALL and
filesort have
disappeared**

This is not always possible

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision WHERE
rev_len > 5500 ORDER BY rev_timestamp ASC\G
***** 1. row *****
      id: 1
    select_type: SIMPLE
      table: revision
        type: range
possible_keys: rev_len_rev_timest
      key: rev_len_rev_timest
    key_len: 5
       ref: NULL
      rows: 38744
  Extra: Using index condition; Using filesort
1 row in set (0.00 sec)
```

The range makes impossible to use the index optimally for the the ORDER BY: either we filter (rev_len) or sort (rev_timestamp)

A Strange Game. The Only Winning Move is Not to Play

```
mysql-5.6.21 (osm) > SELECT * FROM nodes FORCE  
INDEX(version_idx) WHERE version < 15 ORDER BY changeset_id;  
/* type: range, Using filesort */  
2859673 rows in set (30.58 sec)
```

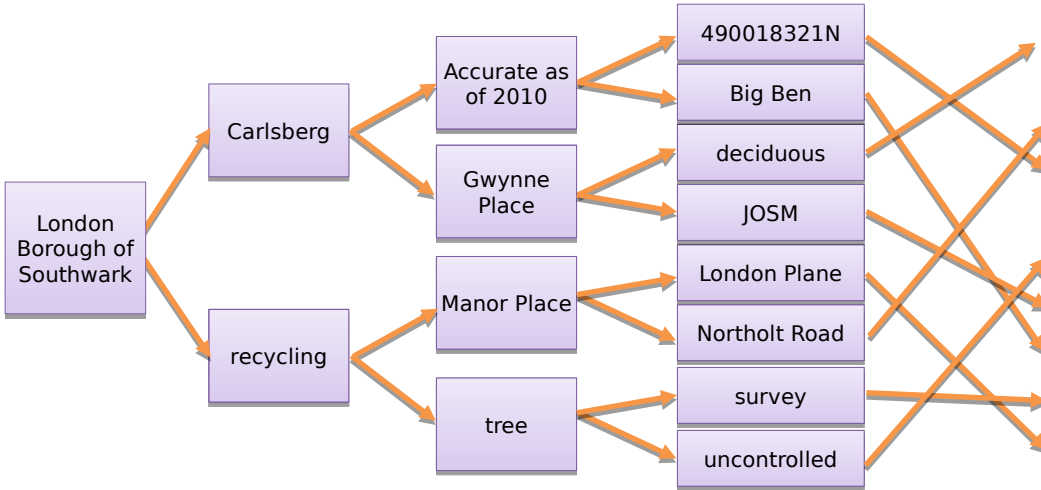
```
mysql-5.6.21 (osm) > SELECT * FROM nodes FORCE  
INDEX(changeset_id_idx) WHERE version < 15 ORDER BY  
changeset_id;  
/* type: index */  
2859673 rows in set (30.92 sec)
```

```
mysql-5.6.21 (osm) > SELECT * FROM nodes WHERE version < 15  
ORDER BY changeset_id;  
/* type: ALL, Using filesort */  
2859673 rows in set (16.54 sec)
```

MyISAM Internals

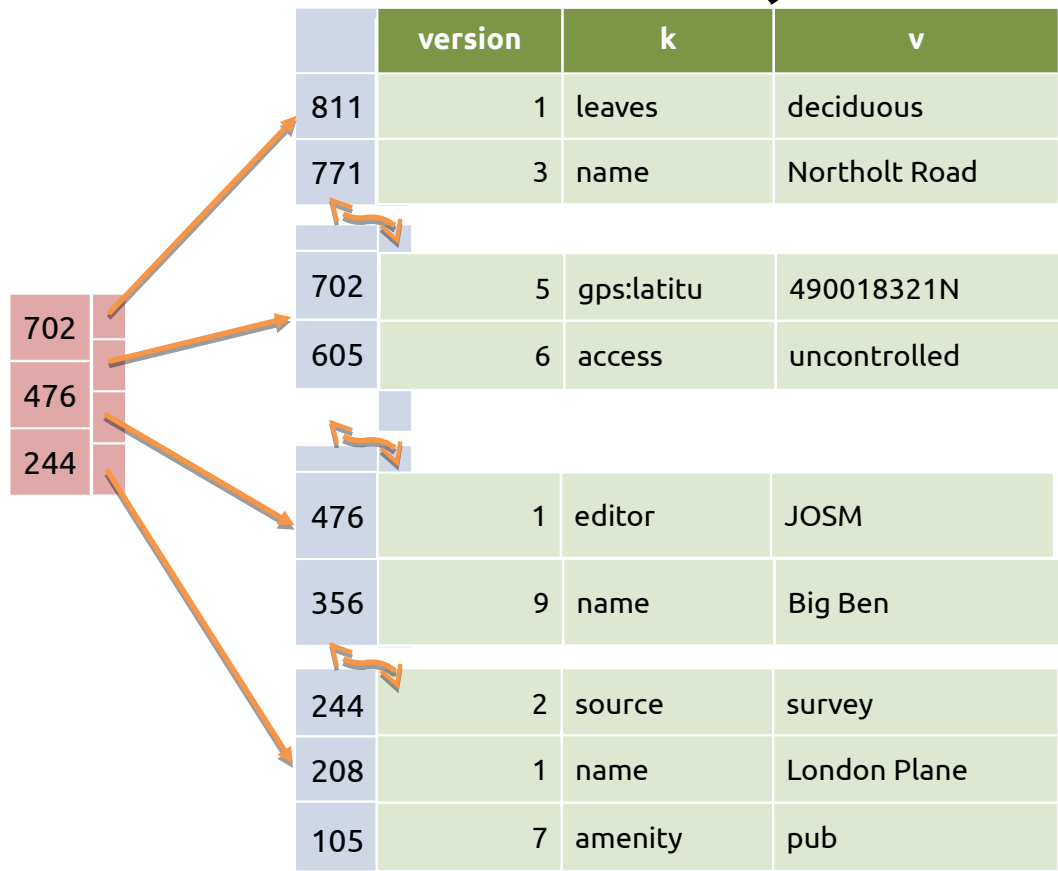
Index (part of revision.MYI)

Data (revision.MYD)



node_id	versio	k	v
234234344545	1	leaves	deciduous
(empty row)			
234234344548	5	gps:latitude	490018321N
234234344549	6	access	uncontrolled
(empty row)			
234234344551	9	name	Big Ben
234234344552	1	source	survey
234234344557	1	name	London Plane
234234344552	2	source	survey
234234344557	2	name	London Plane

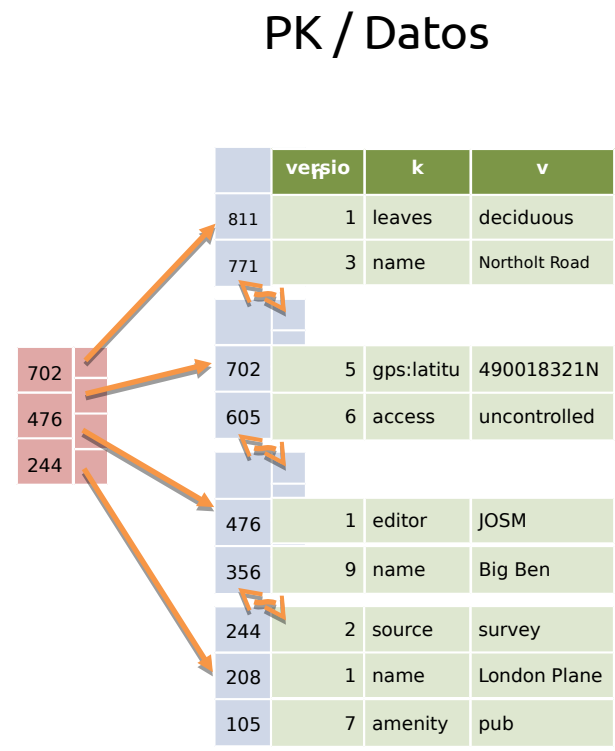

InnoDB Internals (PRIMARY)



**Data
clustered
always
using the
primary
key**

InnoDB Internals (Secondary)

Secondary indexes contain the primary key value



Consequences of using InnoDB (I)

- Every table should have a primary key
 - If one is not defined, MySQL will choose an available NOT NULL unique key
 - If that is not possible, an internal 6-byte row identifier will be generated (not user-accessible)

Consequences of using InnoDB (II)

- Inserting in primary key order is much faster
 - Less fragmentation/page-split
 - Usage of "batch" mode, improving insert speed
- Using auto-increment keys as primary keys can be a good idea for InnoDB

Consequences of using InnoDB (III)

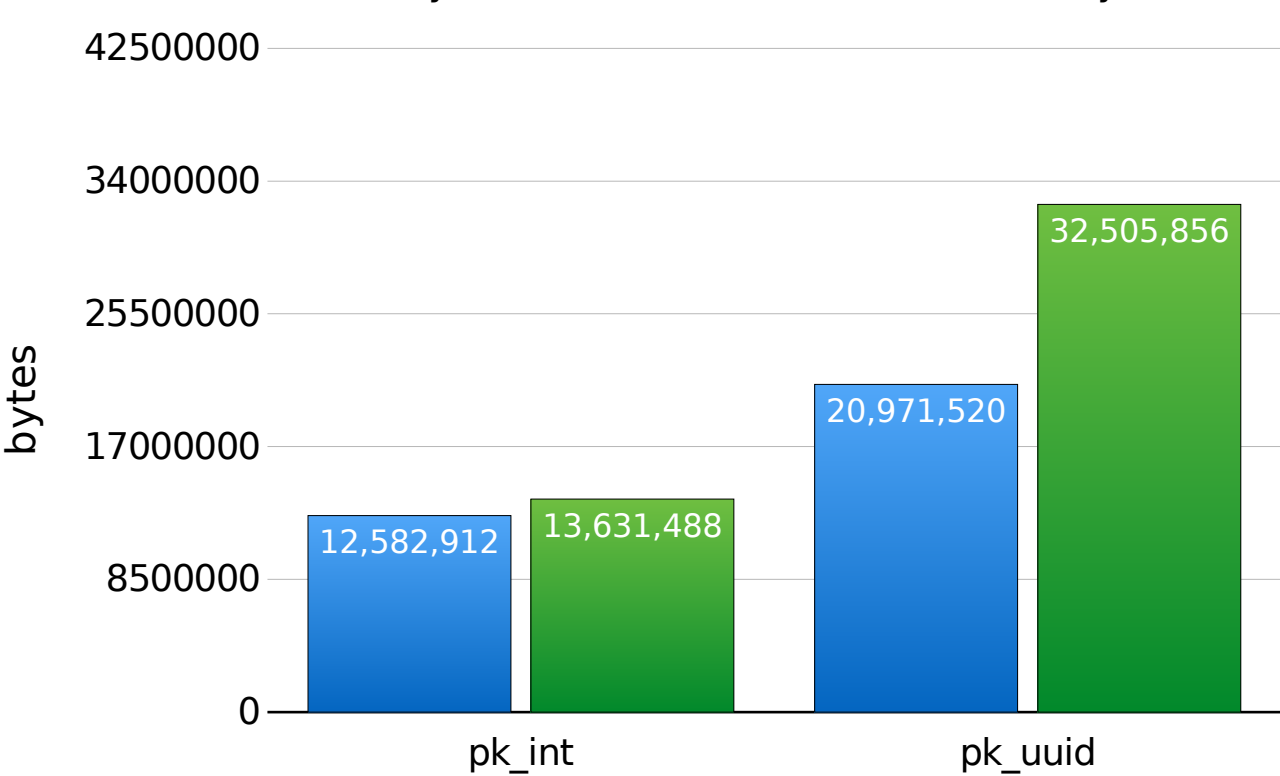
- A very long primary key may increment substantially the size of secondary keys
 - Int or bigint types are recommended instead of UUIDs or other long strings

Differences in size

```
mysql-5.6.21 (osm) > CREATE
TABLE pk_int (id int
PRIMARY KEY auto_increment,
a int,
b int,
c int,
d int);
Query OK, 0 rows affected
(0.16 sec)
```

```
mysql-5.6.21 (osm) > CREATE
TABLE pk_uuid (id char(36)
PRIMARY KEY,
a int,
b int,
c int,
d int);
Query OK, 0 rows affected
(0.04 sec)
```

■ .ibd size (no secondary indexes) ■ .ibd size (with secondary indexes)



Extended primary key optimization

- As the primary key is part of all secondary keys, this can be used “for free”:
 - For row filtering (since MySQL 5.6)
 - To return results in primary key order
 - To avoid reading data from the table (covering index)

Extended Primary Key Example

```
mysql-5.6.21 (osm) > EXPLAIN SELECT node_id FROM nodes WHERE  
changeset_id = 24284 and node_id <> 146472942\G  
***** 1. row *****  
      id: 1  
    select_type: SIMPLE  
      table: nodes  
        type: range  
possible_keys: PRIMARY,changeset_id_idx  
          key: changeset_id_idx  
        key_len: 16  
          ref: NULL  
         rows: 50  
      Extra: Using where; Using index  
1 row in set (0.07 sec)
```

Query Optimization: From 0 to 10 (and up to 5.7)

FULLTEXT SEARCH

Fuzzy Search of “gloucester/Gloucester’s/etc”

- “Typical” way to solve this:

```
mysql-5.7.5 (osm) >  
SELECT way_id as id, v  
FROM way_tags  
WHERE v like '%gloucester%';  
425 rows in set (0.46 sec)
```



Too slow

Let's Add an Index

```
mysql-5.7.5 (osm) > ALTER TABLE way_tags ADD INDEX(v);
Query OK, 0 rows affected (6.44 sec)
Records: 0 Duplicates: 0 Warnings: 0
```

```
mysql-5.7.5 (osm) > SELECT ...;  
425 rows in set (0.38 sec)
```

Still slow, why?

```
mysql-5.7.5 (osm) > EXPLAIN SELECT way as type, way_id as
id, v FROM way_tags WHERE v like '%gloucester%';
```

[illegible]

Fulltext Index

```
mysql-5.7.5 (osm) > ALTER TABLE way_tags add FULLTEXT index(v);
Query OK, 0 rows affected (3.20 sec)
Records: 0  Duplicates: 0  Warnings: 0
```

```
mysql-5.7.5 (osm) > SELECT ...;
425 rows in set (0.00 sec)
```

```
mysql-5.7.5 (osm) > EXPLAIN SELECT 'way' as type, way_id as id, v
FROM way_tags WHERE MATCH(v) AGAINST ('+gloucester*' IN BOOLEAN
MODE);
```

id	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	fulltext	v	v	0	const	1	100.00	Using where; Ft_hints: no_ranking

1 row in set, 1 warning (0.00 sec)

Newer Fulltext Optimizations

```
mysql-5.5.40 (osm) > EXPLAIN SELECT count(*) FROM way_tags_myisam WHERE MATCH(v) AGAINST('gloucester');
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table           | type      | possible_keys | key  | key_len | ref | rows | Extra           |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1  | SIMPLE      | way_tags_myisam | fulltext | v              | v    | 0        |    | 1    | Using where     |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
1 row in set (0.00 sec)
```

```
mysql-5.5.40 (osm) > SHOW STATUS like 'Hand%';
+-----+-----+
| Variable_name | Value |
+-----+-----+
| Handler_commit | 0     |
| Handler_delete | 0     |
...
| Handler_read_first | 0     |
| Handler_read_key   | 0     |
| Handler_read_last  | 0     |
| Handler_read_next  | 425   |
| Handler_read_prev  | 0     |
| Handler_read_rnd   | 0     |
| Handler_read_rnd_next | 0     |
...
| Handler_update     | 0     |
| Handler_write      | 0     |
+-----+-----+
16 rows in set (0.00 sec)
```

Newer Fulltext Optimizations (cont.)

```
mysql-5.7.5 (osm) > EXPLAIN SELECT count(*) FROM way_tags WHERE MATCH(v) AGAINST('gloucester');
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	Select tables optimized away

1 row in set, 1 warning (0.00 sec)

```
mysql-5.7.5 (osm) > SHOW STATUS like 'Hand%';
```

Variable_name	Value
Handler_commit	1
Handler_delete	0
...	
Handler_read_first	0
Handler_read_key	0
Handler_read_last	0
Handler_read_next	0
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_next	0
...	
Handler_update	0
Handler_write	0

18 rows in set (0.00 sec)

It's counting directly from the FULLTEXT index

Open Issues and Limitations

- No postfix support (wildcards)
- Simple Ranking (and different from MyISAM)
- No stemming support
- Some multi-language limitations

More on FULLTEXT InnoDB support:

<http://www.drdobbs.com/database/full-text-search-with-innodb/231902587>

Alternatives

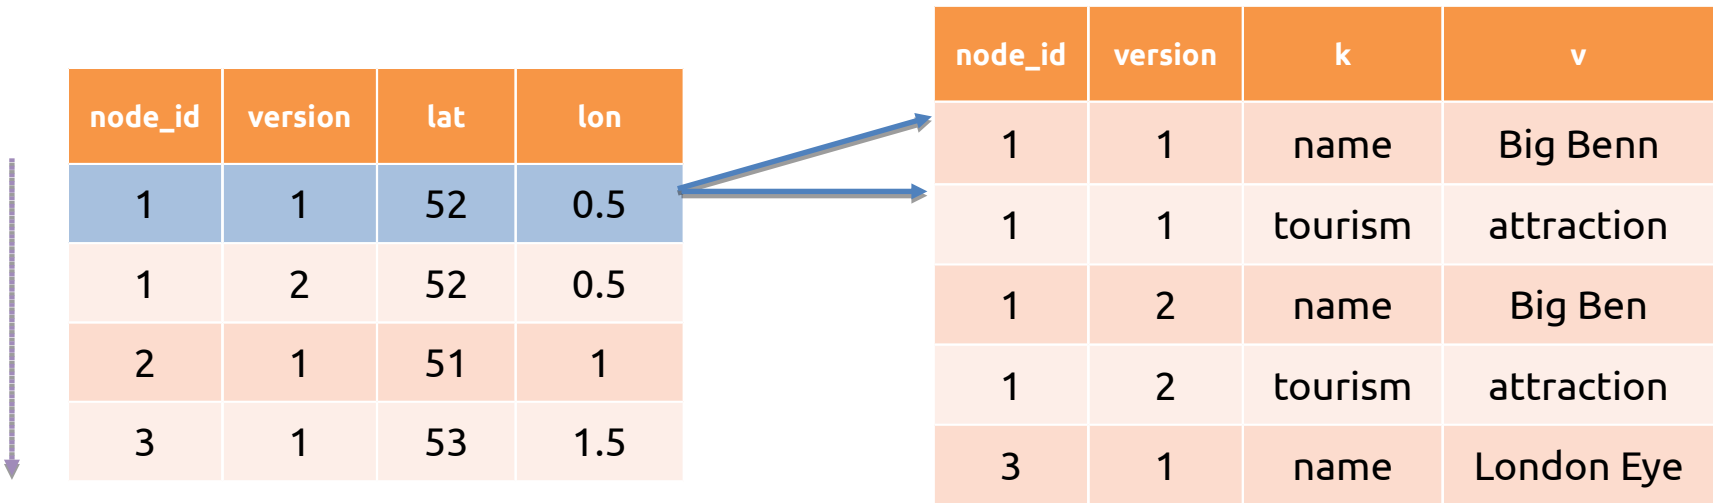
- Apache Lucene
 - Solr
 - Elasticsearch
- Sphinx
 - SphinxSE

Query Optimization: From 0 to 10 (and up to 5.7)

JOINS

(Block) Nested Loop Join

- Until MySQL 5.5 there was only one algorithm to execute a JOIN:



Extra Access type: eq_ref

```
mysql-5.6.21 (osm) > EXPLAIN SELECT * FROM nodes JOIN node_tags
                      USING(node_id, version) WHERE node_tags.v= 'Big Ben'\G
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: node_tags
         type: ref
possible_keys: PRIMARY,v_idx
          key: v_idx
       key_len: 767
         ref: const
        rows: 1
      Extra: Using where; Using index
***** 2. row *****
      id: 1
  select_type: SIMPLE
        table: nodes
         type: eq_ref
possible_keys: PRIMARY,version_idx
          key: PRIMARY
       key_len: 16
         ref: osm.node_tags.node_id,osm.node_tags.version
        rows: 1
      Extra: NULL
2 rows in set (0.00 sec)
```

eq_ref is similar to ref, but allows faster JOINS because, by using a unique key, it only has to search one row for each previous result

JOIN Optimization

- Two main goals:
 - Perform an effective filtering on each table access, if possible using indexes
 - Perform the access in the most efficient table order
- When joining 3 or more tables in a star schema, the "covering index" strategy can have a huge impact

Example: optimize this JOIN (I)

```
SELECT n.node_id, n.latitude, n.longitude
  FROM way_nodes w_n
  JOIN way_tags w_t
    ON w_n.way_id = w_t.way_id and
       w_n.version = w_t.version
  JOIN nodes n
    ON w_n.node_id = n.node_id
  JOIN node_tags n_t
    ON n.node_id = n_t.node_id and
       n.version = n_t.version
 WHERE w_t.k = 'building' and
       n_t.k = 'entrance' and
       n_t.v = 'main';
```



**We start without
secondary indexes**

Example: optimize this JOIN (II)

```

***** 1. row *****
      id: 1
select_type: SIMPLE
      table: w_t
      type: index
possible_keys: PRIMARY
      key: PRIMARY
      key_len: 783
      ref: NULL
      rows: 1335702
Extra: Using where; Using index
***** 2. row *****
      id: 1
select_type: SIMPLE
      table: w_n
      type: ref
possible_keys: PRIMARY
      key: PRIMARY
      key_len: 16
      ref: osm.w_t.way_id,osm.w_t.version
      rows: 3
Extra: NULL

***** 3. row *****
      id: 1
select_type: SIMPLE
      table: n_t
      type: ref
possible_keys: PRIMARY
      key: PRIMARY
      key_len: 8
      ref: osm.w_n.node_id
      rows: 1
Extra: Using where
***** 4. row *****
      id: 1
select_type: SIMPLE
      table: n
      type: eq_ref
possible_keys: PRIMARY
      key: PRIMARY
      key_len: 16
      ref: osm.w_n.node_id,osm.n_t.version
      rows: 1
Extra: Using index
4 rows in set (0.01 sec)

```

mysql-5.6.21 (osm) > SELECT ...
858 rows in set (9.00 sec)

Example: optimize this JOIN (III)

```
mysql-5.6.21 (osm) > ALTER TABLE way_tags ADD  
INDEX k_idx(k);  
Query OK, 0 rows affected (4.80 sec)  
Records: 0 Duplicates: 0 Warnings: 0
```



Creating an index
on way_tags.k

Example: optimize this JOIN (IV)

***** 1. row *****

```
id: 1
select_type: SIMPLE
table: w_t
type: ref
possible_keys: PRIMARY,k_idx
key: k_idx
key_len: 767
ref: const
rows: 452274
Extra: Using where; Using index
```

***** 2. row *****

```
id: 1
select_type: SIMPLE
table: w_n
type: ref
possible_keys: PRIMARY
key: PRIMARY
key_len: 16
ref: osm.w_t.way_id,osm.w_t.version
rows: 3
Extra: NULL
```

mysql-5.6.21 (osm) > SELECT ...
858 rows in set (8.58 sec)

***** 3. row *****

```
id: 1
select_type: SIMPLE
table: n_t
type: ref
possible_keys: PRIMARY
key: PRIMARY
key_len: 8
ref: osm.w_n.node_id
rows: 1
Extra: Using where
```

***** 4. row *****

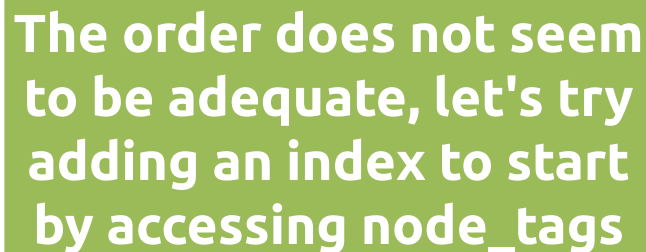
```
id: 1
select_type: SIMPLE
table: n
type: eq_ref
possible_keys: PRIMARY
key: PRIMARY
key_len: 16
ref: osm.w_n.node_id,osm.n_t.version
rows: 1
Extra: NULL
```

4 rows in set (0.00 sec)

**It seems like the index
is not very useful**

Example: optimize this JOIN (V)

```
mysql-5.6.21 (osm) > ALTER TABLE node_tags ADD  
INDEX k_idx(k);  
Query OK, 0 rows affected (2.82 sec)  
Records: 0 Duplicates: 0 Warnings: 0
```



The order does not seem to be adequate, let's try adding an index to start by accessing node_tags

Example: optimize this JOIN (VI)

```

***** 1. row *****
      id: 1
    select_type: SIMPLE
      table: w_t
      type: ref
possible_keys: PRIMARY,k_idx
      key: k_idx
     key_len: 767
       ref: const
      rows: 452274
    Extra: Using where; Using index
***** 2. row *****
      id: 1
    select_type: SIMPLE
      table: w_n
      type: ref
possible_keys: PRIMARY
      key: PRIMARY
     key_len: 16
       ref: osm.w_t.way_id,osm.w_t.version
      rows: 3
    Extra: NULL

***** 3. row *****
      id: 1
    select_type: SIMPLE
      table: n_t
      type: ref
possible_keys: PRIMARY,k_idx
      key: PRIMARY
     key_len: 8
       ref: osm.w_n.node_id
      rows: 1
    Extra: Using where
***** 4. row *****
      id: 1
    select_type: SIMPLE
      table: n
      type: eq_ref
possible_keys: PRIMARY
      key: PRIMARY
     key_len: 16
       ref: osm.w_n.node_id,osm.n_t.version
      rows: 1
    Extra: NULL
4 rows in set (0.00 sec)

```

It keeps using the wrong order, even if we delete the w_t.k_idx index

mysql-5.6.21 (osm) > SELECT ...
 858 rows in set (7.33 sec)

Example: optimize this JOIN (VII)

```
SELECT STRAIGHT_JOIN n.node_id, n.latitude, n.longitude
  FROM node_tags n_t
  JOIN nodes n
    ON n.node_id = n_t.node_id and
       n.version = n_t.version
  JOIN way_nodes w_n
    ON w_n.node_id = n.node_id
  JOIN way_tags w_t
    ON w_n.way_id = w_t.way_id and
       w_n.version = w_t.version
 WHERE w_t.k = 'building' and
        n_t.k = 'entrance' and
        n_t.v = 'main';
```



Let's see why rewriting it
into this query

Example: optimize this JOIN (VIII)

```
***** 1. row *****
  id: 1
  select_type: SIMPLE
  table: n_t
  type: ref
  possible_keys: PRIMARY,k_idx
  key: k_idx
  key_len: 767
  ref: const
  rows: 2390
  Extra: Using index condition; Using where
***** 2. row *****
  id: 1
  select_type: SIMPLE
  table: n
  type: eq_ref
  possible_keys: PRIMARY
  key: PRIMARY
  key_len: 16
  ref: osm.n_t.node_id,osm.n_t.version
  rows: 1
  Extra: NULL
***** 3. row *****
  id: 1
  select_type: SIMPLE
  table: w_n
  type: ALL
  possible_keys: PRIMARY
  key: NULL
  key_len: NULL
  ref: NULL
  rows: 3597858
  Extra: Using where; Using join buffer (Block
Nested Loop)
***** 4. row *****
  id: 1
  select_type: SIMPLE
  table: w_t
  type: eq_ref
  possible_keys: PRIMARY
  key: PRIMARY
  key_len: 783
  ref: osm.w_n.way_id,osm.w_n.version,const
  rows: 1
  Extra: Using where; Using index
4 rows in set (0.00 sec)
```

There is no index on w_n that would allow efficient access

Example: optimize this JOIN (IX)

```
mysql-5.6.21 (osm) > ALTER TABLE way_nodes ADD  
INDEX node_id_idx(node_id);  
Query OK, 0 rows affected (17.77 sec)  
Records: 0 Duplicates: 0 Warnings: 0
```

Example: optimize this JOIN (X)

***** 1. row *****

```
id: 1
select_type: SIMPLE
table: n_t
type: ref
possible_keys: PRIMARY,k_idx
key: k_idx
key_len: 767
ref: const
rows: 2390
```

Extra: Using index condition; Using where

***** 2. row *****

```
id: 1
select_type: SIMPLE
table: n
type: eq_ref
possible_keys: PRIMARY
key: PRIMARY
key_len: 16
ref: osm.n_t.node_id,osm.n_t.version
rows: 1
Extra: NULL
```

Now it starts by
the right table
(without
STRAIGHT_JOIN)

***** 3. row *****

```
id: 1
select_type: SIMPLE
table: w_n
type: ref
possible_keys: PRIMARY,node_id_idx
key: node_id_idx
key_len: 8
ref: osm.n_t.node_id
rows: 1
```

Extra: Using index

***** 4. row *****

```
id: 1
select_type: SIMPLE
table: w_t
type: eq_ref
possible_keys: PRIMARY
key: PRIMARY
key_len: 783
ref: osm.w_n.way_id,osm.w_n.version,const
rows: 1
```

Extra: Using where; Using index

4 rows in set (0.04 sec)

mysql-5.6.21 (osm) > SELECT ...
858 rows in set (0.73 sec)

Example: optimize this JOIN (XI)

```

***** 1. row *****
  id: 1
  select_type: SIMPLE
  table: n_t
  type: ref
  possible_keys: PRIMARY,k_idx,k_v_idx
  key: k_v_idx
  key_len: 1534
  ref: const,const
  rows: 900
  Extra: Using where; Using index
***** 2. row *****
  id: 1
  select_type: SIMPLE
  table: n
  type: eq_ref
  possible_keys: PRIMARY
  key: PRIMARY
  key_len: 16
  ref: osm.n_t.node_id,osm.n_t.version
  rows: 1
  Extra: NULL

```

An index on (k,v) is even better

```

***** 3. row *****
  id: 1
  select_type: SIMPLE
  table: w_n
  type: ref
  possible_keys: PRIMARY,node_id_idx
  key: node_id_idx
  key_len: 8
  ref: osm.n_t.node_id
  rows: 1
  Extra: Using index
***** 4. row *****
  id: 1
  select_type: SIMPLE
  table: w_t
  type: eq_ref
  possible_keys: PRIMARY
  key: PRIMARY
  key_len: 783
  ref: osm.w_n.way_id,osm.w_n.version,const
  rows: 1
  Extra: Using where; Using index
4 rows in set (0.00 sec)

```

mysql-5.6.21 (osm) > SELECT ...
 858 rows in set (0.02 sec)

“New” JOIN methods

- MySQL 5.6 added:
 - Batch Key Access
- MariaDB has since 5.3:
 - Batch Key Access
 - Hash Joins
 - Slightly modified versions of the above ones (with “incremental” buffers to join 3 or more tables)

Multi-range read

- This optimization orders results obtained from a secondary key in primary key/physical order before accessing the rows
 - It may help execution time of queries when disk-bound
 - It requires tuning of the `read_rnd_buffer_size` (size of the buffer used for ordering the results)
- BKA JOINS are based on the mrr optimization

MRR Example (I)

```
mysql-5.6.21 (osm) > EXPLAIN SELECT * FROM nodes WHERE timestamp >=
'2013-07-01 00:00:00' AND timestamp < '2014-01-01 00:00:00'\G
***** 1. row *****
      id: 1
    select_type: SIMPLE
      table: nodes
        type: range
possible_keys: nodes_timestamp_idx
          key: nodes_timestamp_idx
      key_len: 5
         ref: NULL
        rows: 429684
      Extra: Using index condition; Using MRR
1 row in set (0.02 sec)
```

MRR example (II)

[restart]

```
mysql> SET optimizer_switch='mrr=off';
```

```
mysql> SELECT * FROM nodes WHERE timestamp >= '2013-07-01 00:00:00' AND timestamp < '2014-01-01 00:00:00';
```

205617 rows in set (5.16 sec)

```
mysql> SELECT * FROM nodes WHERE timestamp >= '2013-07-01 00:00:00' AND timestamp < '2014-01-01 00:00:00';
```

205617 rows in set (0.60 sec)

[restart]

```
mysql> SET read_rnd_buffer_size=50 * 1024 * 1024;
```

```
mysql> SELECT * FROM nodes WHERE timestamp >= '2013-07-01 00:00:00' AND timestamp < '2014-01-01 00:00:00';
```

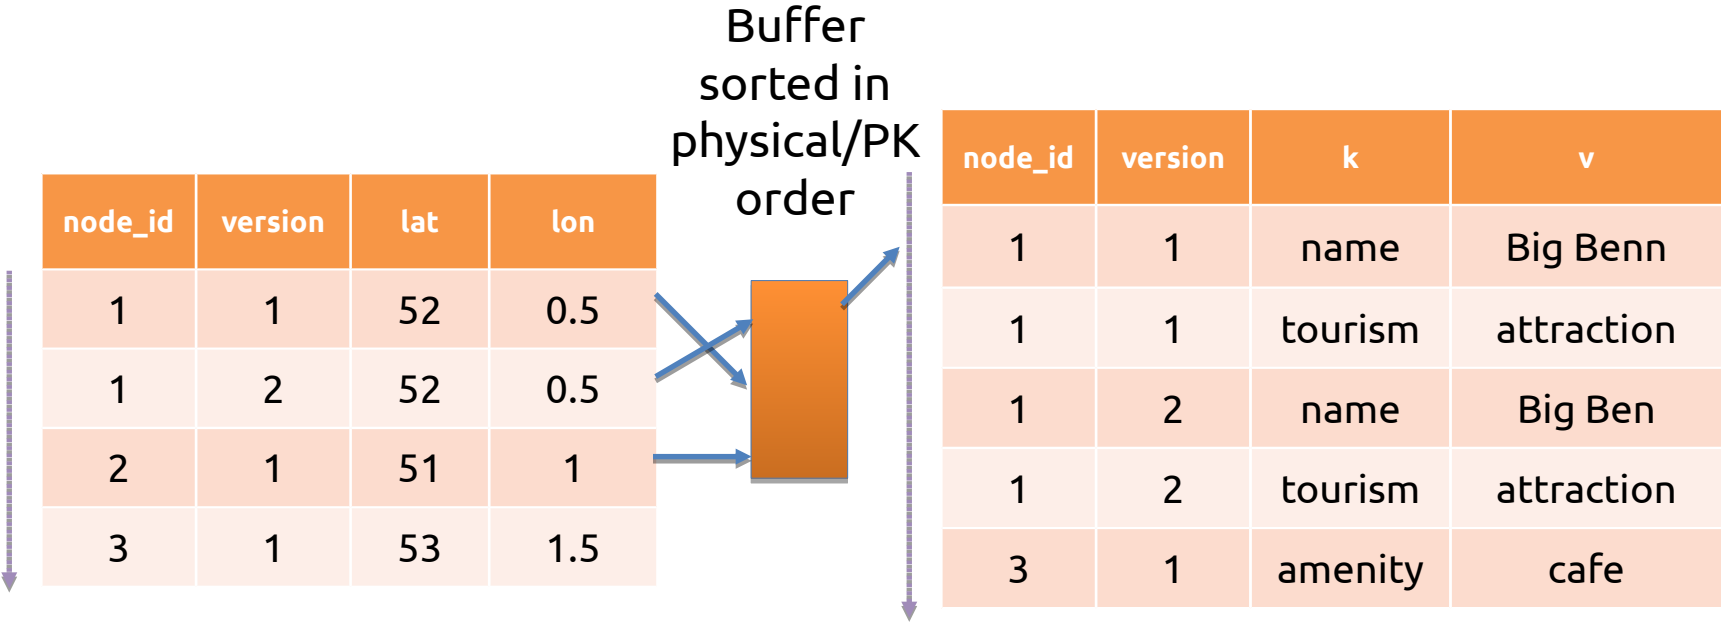
205617 rows in set (2.39 sec)

```
mysql> SELECT * FROM nodes WHERE timestamp >= '2013-07-01 00:00:00' AND timestamp < '2014-01-01 00:00:00';
```

205617 rows in set (0.73 sec)

“Cold” results are significantly better with mrr (but it can impact negatively, too)

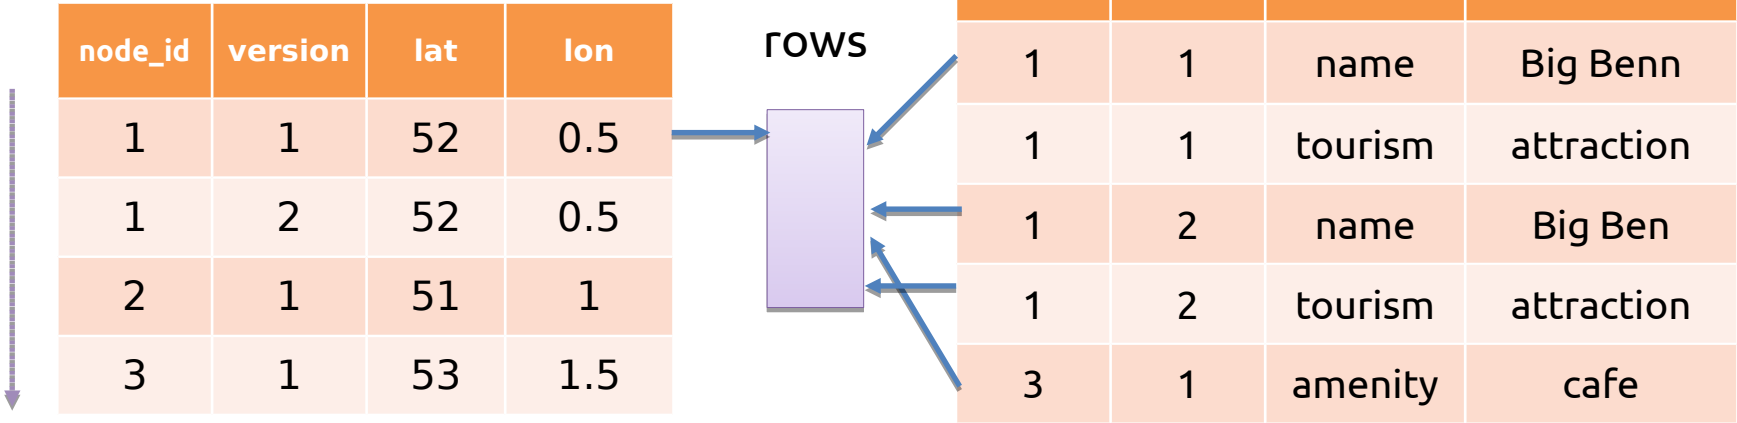
Batch Key Access



Hash Joins

- Only work for equi-joins

Hash table
for faster
access to
rows



MySQL Configuration

- BKA requires changes of default optimizer configuration:

```
mysql-5.7.5 (osm) > SET optimizer_switch= 'mrr=on';
```

```
mysql-5.7.5 (osm) > SET optimizer_switch= 'mrr_cost_based=off';
```

```
mysql-5.7.5 (osm) > SET optimizer_switch= 'batch_key_access=on';
```

- Additionally, configuring the `join_buffer_size` adequately

MariaDB configuration

```
mariadb-10.0.14 (osm) > SET optimizer_switch = 'join_cache_incremental=on';  
mariadb-10.0.14 (osm) > SET optimizer_switch = 'join_cache_hashed=on';  
mariadb-10.0.14 (osm) > SET optimizer_switch = 'join_cache_bka=on';
```

- Enabled by default

```
mariadb-10.0.14 (osm) > SET join_cache_level = 3 (for hash joins)  
mariadb-10.0.14 (osm) > SET join_cache_level = 5 (for BKA)
```

- Also, configure `join_buffer_size` appropriately.
- Hash joins, like BKA, are highly dependent on disk-bound DBs to be effective due to the extra overhead

Nested Loop Join (cold buffers buffer_pool=100MB, join_buffer=4M)

```
mariadb-10.0.14 (osm) > EXPLAIN SELECT
changeset_id, count(*) FROM changesets JOIN nodes
on changesets.id = nodes.changeset_id GROUP BY
visible\G
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: changesets
        type: index
possible_keys: PRIMARY
         key: PRIMARY
      key_len: 8
         ref: NULL
        rows: 69115
  Extra: Using index; Using temporary; Using
filesort
```

```
***** 2. row *****
      id: 1
  select_type: SIMPLE
        table: nodes
        type: ref
possible_keys: changeset_id
         key: changeset_id
      key_len: 8
         ref: osm.changesets.id
        rows: 19
  Extra:
2 rows in set (0.00 sec)

mariadb-10.0.14 (osm) > SELECT visible, count(*)
FROM changesets JOIN nodes on changesets.id =
nodes.changeset_id GROUP BY visible;
+-----+-----+
| visible | count(*) |
+-----+-----+
|      1 | 2865312 |
+-----+-----+
1 row in set (32.86 sec)
```

Hash Join (cold buffers, buffer_pool=100M, join_buffer=4M)

```
mariadb-10.0.14 (osm) > EXPLAIN SELECT
changeset_id, count(*) FROM changesets JOIN
nodes on changesets.id = nodes.changeset_id
GROUP BY visible\G
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: changesets
        type: index
possible_keys: PRIMARY
         key: PRIMARY
      key_len: 8
         ref: NULL
        rows: 69115
  Extra: Using index; Using temporary;
Using filesort
```

```
***** 2. row *****
      id: 1
  select_type: SIMPLE
        table: nodes
        type: hash_ALL
possible_keys: changeset_id
         key: #hash#changeset_id
      key_len: 8
         ref: osm.changesets.id
        rows: 2781732
  Extra: Using join buffer (flat, BNLH join)
2 rows in set (0.00 sec)

mariadb-10.0.14 (osm) > SELECT visible, count(*)
FROM changesets JOIN nodes on changesets.id =
nodes.changeset_id GROUP BY visible;
+-----+-----+
| visible | count(*) |
+-----+-----+
|      1 | 2865312 |
+-----+-----+
1 row in set (6.66 sec)
```

Query Optimization: From 0 to 10 (and up to 5.7)

SUBQUERIES

Access types: unique_subquery/index_subquery

```
mysql-5.6.21 (osm) > EXPLAIN SELECT ***** 2. row *****
* FROM node_tags WHERE v = 'Big Ben'           id: 2
and node_id NOT IN (SELECT node_id             select_type: DEPENDENT SUBQUERY
FROM nodes WHERE tile < 100000000)\G           table: nodes
***** 1. row *****                        type: index_subquery
      id: 1                                possible_keys: PRIMARY,nodes_tile_idx
      select_type: PRIMARY                  key: PRIMARY
      table: node_tags                     key_len: 8
      type: ref                             ref: func
possible_keys: v_idx                       rows: 1
      key: v_idx                           Extra: Using where
      key_len: 767                         2 rows in set (0.00 sec)
      ref: const
      rows: 1
      Extra: Using where; Using
index
```

Unique
subquery
is similar,
but using a
unique or
primary
key

Subqueries in MySQL

- MySQL versions traditionally had very bad press regarding subqueries
 - It was common to recommend rewriting them (when possible) into JOINS
- Since MySQL 5.6, its query execution plans have improved significantly

Lazy Materialization of derived tables

- Option available since MySQL 5.6
 - Improves the execution time of EXPLAIN (it no longer needs to execute subqueries)
 - Derived tables can be indexed automatically at execution time to improve its performance

Derived Table Example

```
mysql-5.5.40 (osm) > EXPLAIN SELECT count(*) FROM (SELECT * FROM nodes WHERE VISIBLE = 1) n JOIN changesets ON n.changeset_id = changesets.id;
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	PRIMARY	<derived2>	ALL	NULL	NULL	NULL	NULL	2865312	
1	PRIMARY	changesets	eq_ref	PRIMARY	PRIMARY	8	n.changeset_id	1	Using index
2	DERIVED	nodes	ALL	NULL	NULL	NULL	NULL	2865521	Using where

3 rows in set (1.42 sec)

```
mysql-5.6.21 (osm) > EXPLAIN SELECT count(*) FROM (SELECT * FROM nodes WHERE VISIBLE = 1) n JOIN changesets ON n.changeset_id = changesets.id;
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	PRIMARY	changesets	index	PRIMARY	PRIMARY	8	NULL	70917	Using index
1	PRIMARY	<derived2>	ref	<auto_key0>	<auto_key0>	8	osm.changesets.id	40	NULL
2	DERIVED	nodes	ALL	NULL	NULL	NULL	NULL	2853846	Using where

3 rows in set (0.00 sec)

Subquery is not executed

Auto-generated index

A Common 5.5 Performance Problem

```
mysql-5.5.40 (osm) > EXPLAIN SELECT * FROM nodes
                        WHERE nodes.changeset_id IN (
                        SELECT changesets.id
                        FROM changesets
                        JOIN users
                        ON changesets.user_id = users.id and users.display_name = 'Steve');
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	PRIMARY	nodes	ALL	NULL	NULL	NULL	NULL	2865521	Using where
2	DEPENDENT SUBQUERY	users	const	PRIMARY,users_display_name_idx	users_display_name_idx	767	const	1	Using index
2	DEPENDENT SUBQUERY	changesets	eq_ref		PRIMARY	8	func	1	Using where

3 rows in set (0.00 sec)

This means that the subquery is executed almost 3 million times

```
mysql-5.5.40 (osm) > SELECT ...;
```

node_id	latitude	longitude	changeset_id	visi
99890	515276425	-1497621	552	
109174	515364532	-1457329	1875	
276538	515324296	-2094688	810	1
442987	515449207	-1275650	1941	1
442988	515449741	-1272860	1941	1
498803	515438432	-1269436	2171	1
138212838	513010180	-1699929	7757299	1

7 rows in set (2.60 sec)

Semijoin Optimization

- The only way to execute certain IN subqueries was to execute them with poor strategy
 - This forced to rewrite certain queries into JOINS or scalar subqueries, when possible
- There are now several additional automatic options:
 - Convert to a JOIN
 - Materialization (including index creation)
 - FirstMatch
 - LooseScan
 - Duplicate Weedout

The Previous Query is Not a Problem in 5.6/5.7/MariaDB 5.3+

```
mysql-5.6.21 (osm) > EXPLAIN SELECT * FROM nodes
                        WHERE nodes.changeset_id IN (
                            SELECT changesets.id
                            FROM change sets
                            JOIN users
                            ON changesets.user_id = users.id and users.display_name = 'Steve');
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	users	const	PRIMARY,users_display_name_idx	users_display_name_idx	767	const	1	Using index
1	SIMPLE	changesets	ALL	PRIMARY	NULL	NULL	NULL	70917	Using where
1	SIMPLE	nodes	ref	changeset_id	changeset_id	8	osm.changesets.id	21	NULL

3 rows in set (0.00 sec)

```
mysql-5.6.21 (osm) > SELECT ...;
```

node_id	latitude	longitude	changeset_id	visible	timestamp	tile	version
99890	515276425	-1497621	552	1	2005-10-25 00:35:24	2062268512	1
...							
138212838	513010180	-1699929	7757299	1	2011-04-03 18:14:14	2062220563	6

7 rows in set (0.02 sec)



First Match Strategy

```
mysql-5.7.5 (osm) > EXPLAIN SELECT * FROM changesets WHERE id IN (SELECT changeset_id FROM nodes)\G
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: changesets
  partitions: NULL
         type: ALL
possible_keys: PRIMARY
          key: NULL
        key_len: NULL
         ref: NULL
         rows: 70917
   filtered: 100.00
      Extra: NULL
***** 2. row *****
      id: 1
  select_type: SIMPLE
        table: nodes
  partitions: NULL
         type: ref
possible_keys: changeset_id
          key: changeset_id
        key_len: 8
         ref: osm.changesets.id
         rows: 33
   filtered: 100.00
      Extra: Using index; FirstMatch(changesets)
2 rows in set, 1 warning (0.00 sec)
```

It is converting the ref into an eq_ref, shot-circuiting the execution

Enabling and disabling materialization, semijoin, etc

```
mysql-5.7.8 (osm) > SHOW VARIABLES like 'optimizer_switch'\G
```

```
***** 1. row *****
```

```
Variable_name: optimizer_switch
```

```
Value:
```

```
index_merge=on,index_merge_union=on,index_merge_sort_union=on,index_merge_in  
tersection=on,engine_condition_pushdown=on,index_condition_pushdown=on,mrr=o  
n,mrr_cost_based=on,block_nested_loop=on,batched_key_access=off,materializat  
ion=on,semijoin=on,loosescan=on,firstmatch=on,duplicateweedout=on,subquery_m  
aterialization_cost_based=on,use_index_extensions=on,condition_fanout_filter  
=on,derived_merge=on
```

```
1 row in set (0.00 sec)
```

Query Optimization: From 0 to 10 (and up to 5.7)

QUERY PROFILING

Which Queries Should I Optimize First?

- My two favorite methods:
 - pt-query-digest
 - PERFORMANCE_SCHEMA
- I prefer pt-query-digest for long-term reports, PERFORMANCE_SCHEMA for more real-time evaluation and fine-tuning
 - Also, PERFORMANCE_SCHEMA was not “ready” until MySQL 5.6

pt-query-digest

- It is a 3rd party tool written in Perl, originally created by Baron Schwartz
- It requires activation of the slow log:
 - `SET GLOBAL slow_query_log = 1;`
 - `SET long_query_time = 0;`
- In Percona Server and MariaDB it can provide extra information:
 - `SHOW GLOBAL VARIABLES like 'log_slow_verbosity';`



Be careful
with extra
IO and
latency!

pt-query-digest Execution (I)

```
# 1094.7s user time, 9.4s system time, 141.22M rss, 205.21M vsz
# Current date: Wed Jul 1 07:32:28 2015
# Hostname: db1018
# Files: STDIN
# Overall: 4.66M total, 640 unique, 53.47 QPS, 0.02x concurrency _____
# Time range: 2015-06-30 07:00:10 to 2015-07-01 07:11:37
# Attribute          total      min      max      avg      95%    stddev  median
# =====
# Exec time          1320s      1us      3s      283us    332us     3ms    152us
# Lock time           238s        0     13ms     51us     93us    39us     52us
# Rows sent           5.02M        0    4.16k     1.13     1.96     8.69     0.99
# Rows examine        9.50M        0 232.93k     2.14     3.89   261.15     0.99
# Merge passes         0         0         0         0         0         0         0
# Query size          1.06G       17  67.89k    243.89   511.45   368.99   192.76
# Boolean:
# Filesort           8% yes,   91% no
# Full scan          94% yes,    5% no
# Priority que        3% yes,   96% no
# Tmp table          29% yes,   70% no
# Tmp table on       1% yes,   98% no
```

Actual execution
on Wikipedia
production
servers

pt-query-digest Execution (II)

# Profile									
#	Rank	Query ID	Response time		Calls	R/Call	V/M	Item	
#	====	=====	=====		=====	=====	=====	=====	
#	1	0xSANITIZED	242.2765	18.4%	691005	0.0004	0.00	SELECT	revision page user
#	2	0xSANITIZED	204.7052	15.5%	80863	0.0025	0.01	SELECT	revision page user
#	3	0xSANITIZED	162.8476	12.3%	1025179	0.0002	0.00	SELECT	page
#	4	0xSANITIZED	68.1164	5.2%	93928	0.0007	0.01	SELECT	revision page user
#	5	0xSANITIZED	66.8302	5.1%	354562	0.0002	0.00	SELECT	page revision
#	6	0xSANITIZED	57.0374	4.3%	211631	0.0003	0.00	SELECT	page revision
#	7	0xSANITIZED	44.0751	3.3%	6925	0.0064	0.07	SELECT	page categorylinks
category									
#	8	0xSANITIZED	35.0655	2.7%	9689	0.0036	0.00	SELECT	text
#	9	0xSANITIZED	29.4363	2.2%	152259	0.0002	0.00	SELECT	page
#	10	0xSANITIZED	24.1864	1.8%	176927	0.0001	0.00	SELECT	msg_resource
#	11	0xSANITIZED	23.7016	1.8%	144807	0.0002	0.00	SELECT	page_restrictions
#	12	0xSANITIZED	16.6547	1.3%	10135	0.0016	0.03	SELECT	revision
#	13	0xSANITIZED	15.0564	1.1%	263809	0.0001	0.00	SET	

pt-query-digest Execution (III)

```
# Query 1: 7.93 QPS, 0.00x concurrency, ID 0xSANITIZED at byte 1553864032
# This item is included in the report because it matches -limit.
# Scores: V/M = 0.00
# Time range: 2015-06-30 07:00:10 to 2015-07-01 07:11:37
# Attribute      pct  total      min      max      avg      95%    stddev  median
# =====
# Count          14   691005
# Exec time      18    242s    163us    91ms    350us    348us    563us    301us
# Lock time      26     63s     47us     7ms     91us    103us    14us     84us
# Rows sent      12  657.18k      0       1     0.97     0.99     0.16     0.99
# Rows examine   6  657.18k      0       1     0.97     0.99     0.16     0.99
# Query size     31  345.42M    501     749   524.16   537.02     9.22   511.45
# String:
# Databases      itwiki (225976/32%), enwiktiona... (219461/31%)... 15 more
# Hosts
# Users          wikiuser
# Query_time distribution
#   1us
#  10us
# 100us #####
#   1ms #
#  10ms #
# 100ms
#   1s
#  10s+
# Tables
#   SHOW TABLE STATUS FROM `enwiktionary` LIKE 'revision'\G
#   SHOW CREATE TABLE `enwiktionary`.`revision`\G
#   SHOW TABLE STATUS FROM `enwiktionary` LIKE 'page'\G
#   SHOW CREATE TABLE `enwiktionary`.`page`\G
#   SHOW TABLE STATUS FROM `enwiktionary` LIKE 'user'\G
#   SHOW CREATE TABLE `enwiktionary`.`user`\G
# EXPLAIN /*!50100 PARTITIONS*/
SELECT /* Revision::fetchFromConds SANITIZED */ * FROM `revision`
INNER JOIN `page` ON ((page_id = rev_page)) LEFT JOIN `user` ON
((rev_user != 0) AND (user_id = rev_user)) WHERE page_namespace = '0' AND
page_title = 'SANITIZED' AND (rev_id=page_latest) LIMIT 1\G
```

PERFORMANCE_SCHEMA

- Monitoring schema (engine) enabled by default since MySQL 5.6
 - `performance_schema = 1` (it is not dynamic)
- Deprecates the old query profiling
- It is way more user-friendly when combined with the [SYS schema/ps_helper](#) (a set of views and stored procedures created by Mark Leith)
 - Included by default since 5.7.7

Installation of the SYS Schema for 5.6/MariaDB

```
$ git clone https://github.com/MarkLeith/mysql-sys.git
Cloning into 'mysql-sys'...
remote: Counting objects: 926, done.
remote: Compressing objects: 100% (73/73), done.
remote: Total 926 (delta 35), reused 6 (delta 2)
Receiving objects: 100% (926/926), 452.19 KiB | 225.00
KiB/s, done.
Resolving deltas: 100% (584/584), done.
$ cd mysql-sys/
$ ~/sandboxes/msb_5_6_24/use < sys_56.sql
```

Example Usage: Discovering Unused Indexes

```
mysql-5.7.8 (osm) > SELECT * FROM sys.schema_unused_indexes LIMIT 5;
```

object_schema	object_name	index_name
osm	acls	acls_k_idx
osm	changeset_tags	changeset_tags_id_idx
osm	current_nodes	current_nodes_timestamp_idx
osm	current_nodes	current_nodes_tile_idx
osm	current_relations	current_relations_timestamp_idx

5 rows in set (0.04 sec)

```
mysql-5.7.8 (osm) > SELECT * FROM current_nodes WHERE tile = 100;
```

...

```
mysql-5.7.8 (osm) > SELECT * FROM sys.schema_unused_indexes LIMIT 5;
```

object_schema	object_name	index_name
osm	acls	acls_k_idx
osm	changeset_tags	changeset_tags_id_idx
osm	current_nodes	current_nodes_timestamp_idx
osm	current_relations	current_relations_timestamp_idx
osm	current_relations	changeset_id

5 rows in set (0.03 sec)



Example Usage: Slow Queries (ordered by server time)

```
mysql-5.7.8 (osm) > SELECT * FROM sys.statement_analysis LIMIT 10\G
```

```
***** 1. row *****
      query: SELECT `way_id` AS
`id` , `v` FROM `way_tags` WHERE `v` LIKE ?
      db: osm
      full_scan: *
      exec_count: 15
      err_count: 0
      warn_count: 0
      total_latency: 7.83 s
      max_latency: 1.33 s
      avg_latency: 521.84 ms
      lock_latency: 17.94 ms
      rows_sent: 6779
      rows_sent_avg: 452
      rows_examined: 20152155
      rows_examined_avg: 1343477
      rows_affected: 0
      rows_affected_avg: 0
      tmp_tables: 0
      tmp_disk_tables: 0
      rows_sorted: 0
      sort_merge_passes: 0
      digest:
21f90695b1ebf20a5f4d4c1e5e860f58
      first_seen: 2014-11-01 17:04:51
      last_seen: 2014-11-01 17:05:22
```

Example Usage: Top Queries Creating Temporary Tables

```
mysql-5.7.8 (osm) > SELECT * FROM  
sys.statements_with_temp_tables WHERE db = 'osm' LIMIT 10\G
```

```
***** 1. row *****
```

```
query: SELECT ? AS TYPE , `node_id` A ... gs` WHERE `k` = ? AND `v` = ?
```

```
db: osm
```

```
exec_count: 11
```

```
total_latency: 7.57 s
```

```
memory_tmp_tables: 11
```

```
disk_tmp_tables: 0
```

```
avg_tmp_tables_per_query: 1
```

```
tmp_tables_to_disk_pct: 0
```

```
first_seen: 2014-11-01 17:33:55
```

```
last_seen: 2014-11-01 17:34:45
```

```
digest: 5e6e82799b7c7c0e5c57cfe63eb98d5d
```


Example Usage: Top Queries

Creating Temporary Tables (cont.)

```
mysql-5.7.8 (osm) > SELECT DIGEST_TEXT FROM performance_schema.events_statements_summary_by_digest
WHERE digest = '5e6e82799b7c7c0e5c57cfe63eb98d5d'\G
```

```
***** 1. row *****
DIGEST_TEXT: SELECT ? AS TYPE , `node_id` AS `id` FROM `node_tags` WHERE `k` = ? AND `v` = ? UNION
SELECT ? AS TYPE , `way_id` AS `id` FROM `way_tags` WHERE `k` = ? AND `v` = ? UNION SELECT ? AS
TYPE , `relation_id` AS `id` FROM `relation_tags` WHERE `k` = ? AND `v` = ?
1 row in set (0.00 sec)
```

```
mysql-5.7.8 (osm) > EXPLAIN SELECT 'node' as type, node_id as id FROM node_tags WHERE k='amenity'
and v='cafe' UNION SELECT 'way' as type, way_id as id FROM way_tags WHERE k='amenity' and v='cafe'
UNION SELECT 'relation' as type, relation_id as id FROM relation_tags WHERE k='amenity' and
v='cafe';
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	PRIMARY	node_tags	NULL	ALL	NULL	NULL	NULL	NULL	851339	0.00	Using where
2	UNION	way_tags	NULL	ALL	NULL	NULL	NULL	NULL	1331016	0.00	Using where
3	UNION	relation_tags	NULL	ALL	NULL	NULL	NULL	NULL	63201	0.00	Using where
NULL	UNION RESULT	<union1,2,3>	NULL	ALL	NULL	NULL	NULL	NULL	NULL	NULL	Using temporary

4 rows in set, 1 warning (0.01 sec)

Query Optimization: From 0 to 10 (and up to 5.7)

GENERAL OPTIMIZER IMPROVEMENTS

EXPLAIN FOR CONNECTION (SHOW EXPLAIN FOR)

mysql 5.7.8 > SHOW PROCESSLIST;

+-----+-----+-----+-----+-----+-----+-----+-----+							
+-----+-----+-----+-----+-----+-----+-----+-----+							
Id	User	Host	db	Command	Time	State	Info
+-----+-----+-----+-----+-----+-----+-----+-----+							
4	msandbox	localhost	NULL	Query	1	starting	SHOW PROCESSLIST
8	msandbox	localhost	nlwiktory	Query	6	update	INSERT INTO `pageli
(74239,0,0,'WikiWoordenboek:Genus'),(74240,0,0,'WikiWoordenboek:Genus							
+-----+-----+-----+-----+-----+-----+-----+-----+							

2 rows in set (0.14 sec)

mysql 5.7.8 > EXPLAIN FOR CONNECTION 8;

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+											
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+											
id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+											
1	INSERT	pagelinks	NULL	ALL	NULL	NULL	NULL	NULL	NULL	NULL	NULL
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+											


1 row in set (0.03 sec)

Use "SHOW EXPLAIN FOR #" with MariaDB 10

MySQL 5.7 Optimizer Tweaks

```
mysql-5.5.40 (osm) > SELECT *  
                      FROM nodes  
                      JOIN node_tags  
                        ON node_tags.node_id = nodes.node_id  
                      WHERE nodes.latitude  
                         BETWEEN 517000000 and 520000000\G  
  
59 rows in set (1.37 sec)
```

```
mysql-5.5.40 (osm) > SELECT STRAIGHT_JOIN *  
                      FROM nodes  
                      JOIN node_tags  
                        ON node_tags.node_id =  
                      WHERE nodes.latitude  
                         BETWEEN 517000000 and 520000000\G  
  
59 rows in set (0.86 sec)
```



This condition is
very selective,
but there is no
index available

Why does it take the wrong table order?

```
mysql-5.5.40 (osm) > EXPLAIN EXTENDED SELECT * FROM nodes JOIN node_tags ON node_tags.node_id = nodes.node_id WHERE
nodes.latitude BETWEEN 517000000 and 520000000\G
***** 1. row *****
      id: 1
    select_type: SIMPLE
          table: node_tags
          type: ALL
possible_keys: PRIMARY
          key: NULL
         key_len: NULL
          ref: NULL
         rows: 839031
    filtered: 100.00
      Extra:
***** 2. row *****
      id: 1
    select_type: SIMPLE
          table: nodes
          type: ref
possible_keys: PRIMARY
          key: PRIMARY
         key_len: 8
          ref: osm.node_tags.node_id
         rows: 1
    filtered: 100.00
      Extra: Using where
2 rows in set, 1 warning (0.00 sec)
```

Unindexed Columns Are not Accounted

```
mysql-5.5.40 (osm) > EXPLAIN EXTENDED SELECT STRAIGHT_JOIN * FROM nodes JOIN node_tags ON node_tags.node_id = nodes.node_id
WHERE nodes.latitude BETWEEN 517000000 and 520000000\G
***** 1. row *****
    id: 1
  select_type: SIMPLE
    table: nodes
    type: ALL
possible_keys: PRIMARY
    key: NULL
   key_len: NULL
    ref: NULL
    rows: 2865521
  filtered: 100.00
  Extra: Using where
***** 2. row *****
    id: 1
  select_type: SIMPLE
    table: node_tags
    type: ref
possible_keys: PRIMARY
    key: PRIMARY
   key_len: 8
    ref: osm.nodes.node_id
    rows: 1
  filtered: 100.00
  Extra:
2 rows in set, 1 warning (0.00 sec)
```

The optimizer assumes that all rows will be returned

What's new in 5.7?

```
mysql-5.7.5 (osm) > EXPLAIN SELECT * FROM nodes JOIN node_tags ON node_tags.node_id = nodes.node_id WHERE nodes.latitude BETWEEN 517000000 and 5200000000\G
***** 1. row *****
  id: 1
select_type: SIMPLE
  table: nodes
  partitions: NULL
  type: ALL
possible_keys: PRIMARY
  key: NULL
  key_len: NULL
  ref: NULL
  rows: 2773853
  filtered: 11.11
  Extra: Using where
***** 2. row *****
  id: 1
select_type: SIMPLE
  table: node_tags
  partitions: NULL
  type: ref
possible_keys: PRIMARY
  key: PRIMARY
  key_len: 8
  ref: osm.nodes.node_id
  rows: 3
  filtered: 100.00
  Extra: NULL
2 rows in set, 1 warning (0.00 sec)
```

It is actually 0.002%, but it is good enough

If things go wrong -our Friend optimizer_switch

```
mysql-5.7.8 (osm) > SET optimizer_switch='condition_fanout_filter=off';  
Query OK, 0 rows affected (0.00 sec)
```



Improved EXPLAIN FORMAT=JSON

```
"nested_loop": [  
  {  
    "table": {  
      "table_name": "nodes",  
      "access_type": "ALL",  
      "possible_keys": [  
        "PRIMARY"  
      ],  
      "rows_examined_per_scan": 2773853,  
      "rows_produced_per_join": 308175,  
      "filtered": 11.11,  
      "cost_info": {  
        "read_cost": "512783.58",  
        "eval_cost": "61635.01",  
        "prefix_cost": "574418.60",  
        "data_read_per_join": "21M"  
      },  
    },  
  ],  
],
```

**Cost
information
is now
included**

**Engine
statistics are
now floats for
improved
precision**

Optimizer Cost Tuning



Configurable Costs

```
mysql 5.7.8> SELECT * FROM mysql.server_cost;
```

cost_name	cost_value	last_update	comment
disk_temptable_create_cost	NULL	2015-09-20 20:48:10	NULL
disk_temptable_row_cost	NULL	2015-09-20 20:48:10	NULL
key_compare_cost	NULL	2015-09-20 20:48:10	NULL
memory_temptable_create_cost	NULL	2015-09-20 20:48:10	NULL
memory_temptable_row_cost	NULL	2015-09-20 20:48:10	NULL
row_evaluate_cost	NULL	2015-09-20 20:48:10	NULL

6 rows in set (0.00 sec)

The unit is
“read of a
random
data page”

```
mysql 5.7.8> SELECT * FROM mysql.engine_cost;
```

engine_name	device_type	cost_name	cost_value	last_update	comment
default	0	io_block_read_cost	NULL	2015-09-20 20:48:10	NULL
default	0	memory_block_read_cost	NULL	2015-09-20 20:48:10	NULL

2 rows in set (0.00 sec)

Changing Costs Example

```
mysql-5.7.5 (osm) > EXPLAIN FORMAT=JSON SELECT 'node' as type, node_id as id FROM
node_tags WHERE k='amenity' and v='cafe' UNION SELECT 'way' as type, way_id as id FROM
way_tags WHERE k='amenity' and v='cafe' UNION SELECT 'relation' as type, relation_id as
id FROM relation_tags WHERE k='amenity' and v='cafe'\G
```

```
  "cost_info": {
    "query_cost": "22567.80"
```

```
1 row in set, 1 warning (0.00 sec)
```

```
mysql-5.7.5 (osm) > UPDATE mysql.server_cost SET cost_value = 10;
```

```
mysql-5.7.5 (osm) > FLUSH OPTIMIZER_COSTS;
```

```
<session restart>
```

```
mysql-5.7.5 (osm) > EXPLAIN ... \G
```

```
  "cost_info": {
    "query_cost": "661371.00"
```

```
1 row in set, 1 warning (0.00 sec)
```

More info on usage in the manual: <http://dev.mysql.com/doc/refman/5.7/en/cost-model.html>

Still waiting for...

- Utility to analyze the underlying technology (HD, SSD, memory) and filling up the tables automatically
- Buffer “hotness”-aware statistics
- Statistics for non-indexed columns/histograms
 - MariaDB has histograms since 10

Engine-Independent Statistics/Histograms in MariaDB 10

```
mariadb-10.0.14 (osm) > SET histogram_size = 255;  
Query OK, 0 rows affected (0.00 sec)
```

```
mariadb-10.0.14 (osm) > SET use_stat_tables = 2;  
Query OK, 0 rows affected (0.01 sec)
```

```
mariadb-10.0.14 (osm) > ANALYZE TABLE node_tags;  
+-----+-----+-----+-----+  
| Table      | Op      | Msg_type | Msg_text                                     |  
+-----+-----+-----+-----+  
| osm.node_tags | analyze | status   | Engine-independent statistics collected |  
| osm.node_tags | analyze | status   | OK                                       |  
+-----+-----+-----+-----+  
2 rows in set (3.01 sec)
```

```
mariadb-10.0.14 (osm) > ANALYZE TABLE nodes;  
+-----+-----+-----+-----+  
| Table      | Op      | Msg_type | Msg_text                                     |  
+-----+-----+-----+-----+  
| osm.nodes | analyze | status   | Engine-independent statistics collected |  
| osm.nodes | analyze | status   | OK                                       |  
+-----+-----+-----+-----+  
2 rows in set (32.52 sec)
```

```
mariadb-10.0.14 (osm) > SET optimizer_use_condition_selectivity = 3; (or 4 for histograms)  
Query OK, 0 rows affected (0.00 sec)
```

Better Stats!

```
mariadb-10.0.14 (osm) > EXPLAIN EXTENDED SELECT * FROM nodes JOIN node_tags ON node_tags.node_id = nodes.node_id WHERE
nodes.latitude BETWEEN 517000000 and 520000000\G
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: nodes
        type: ALL
possible_keys: PRIMARY
         key: NULL
      key_len: NULL
         ref: NULL
        rows: 2865312
   filtered: 0.39
      Extra: Using where
***** 2. row *****
      id: 1
  select_type: SIMPLE
        table: node_tags
        type: ref
possible_keys: PRIMARY
         key: PRIMARY
      key_len: 8
         ref: osm.nodes.node_id
        rows: 3
   filtered: 100.00
      Extra:
2 rows in set, 1 warning (0.01 sec)
```



Other Changes/Bugs Fixed

- UNION ALL does not create temporary tables, returns tables faster
- (a, b) IN ((1, 2), (2, 3)) can use index ranges
- EXPLAIN EXTENDED is now the default behavior
- I.C. Pushdown support for partitioned tables
- IGNORE clause meaning has been standardized between sentence types
- Increased default for optimizer_search_depth

Query Optimization: From 0 to 10 (and up to 5.7)

COMPUTED/VIRTUAL COLUMNS

Syntax

ALTER TABLE nodes

```
ADD COLUMN lon DECIMAL (10, 7)
  as (longitude/100000000) VIRTUAL,
ADD COLUMN lat DECIMAL (9, 7)
  as (latitude/100000000) VIRTUAL;
```



VIRTUAL is
optional
(default)

- They can be used to simplify SELECTS, calculating values on the fly
- Non accessed rows are not calculated

Functional Indexes

- Before 5.7.8:

```
mysql-5.7.5 (osm) > ALTER TABLE nodes add index(lon);  
ERROR 1951 (HY000): Key/Index cannot be defined on a non-  
stored virtual column.
```


- Now:

```
mysql-5.7.8 (osm) > ALTER TABLE nodes add index(lon);  
Query OK, 0 rows affected (16.54 sec)  
Records: 0 Duplicates: 0 Warnings: 0
```

- This effectively is an implementation of functional indexes, allowing to solve previous query optimization issues without the overhead of an extra column

Do you remember this query?

```
MariaDB [nlwiktionary]> EXPLAIN SELECT * FROM revision WHERE
substr(rev_timestamp, 5, 2) = '09'\G
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: revision
         type: ALL
possible_keys: NULL
          key: NULL
        key_len: NULL
         ref: NULL
         Rows: 173154
      Extra: Using where
1 row in set (0.00 sec)
```



Can you think a way to improve this query?

Now we can solve it like this

```
mysql-5.7.8> ALTER TABLE revision ADD COLUMN rev_month tinyint AS
(substr(rev_timestamp, 5, 2)) VIRTUAL, ADD INDEX rev_month (rev_month);
Query OK, 820308 rows affected (3 min 29.48 sec)
Records: 820308 Duplicates: 0 Warnings: 0
```

```
mysql-5.7.8> EXPLAIN SELECT * FROM revision WHERE rev_month = 9\G
***** 1. row *****
      id: 1
  select_type: SIMPLE
        table: revision
  partitions: NULL
         type: ref
possible_keys: rev_month
          key: rev_month
        key_len: 2
         ref: const
        rows: 104112
   filtered: 100.00
      Extra: NULL
1 row in set, 1 warning (0.01 sec)
```

**The column does not
take space, only the
index**

Stored Columns

```
mysql-5.7.8 (osm) > ALTER TABLE nodes CHANGE lat lat DECIMAL (9, 7) as  
(latitude/100000000) STORED;  
ERROR 1954 (HY000): 'Changing the STORED status' is not supported for virtual  
columns.
```



They dropped the 'yet' :-)

```
mysql-5.7.8 (osm) > ALTER TABLE nodes DROP COLUMN lat, ADD COLUMN lat DECIMAL  
(9, 7) as (latitude/100000000) STORED;  
Query OK, 2865312 rows affected (4 min 51.05 sec)  
Records: 2865312 Duplicates: 0 Warnings: 0
```



MariaDB uses
the
'PERMANENT'
keyword

Features and Limitations

- ~~Virtual “non-stored” columns cannot be indexed~~
- “Stored columns” can be PRIMARY, UNIQUE, FOREIGN and MULTI keys
- It cannot contain subqueries or other tables or rows references
- It cannot contain used-defined stored functions
- It can contain server and user-defined variables
- It can be computed based on other virtual columns

Query Optimization: From 0 to 10 (and up to 5.7)

QUERY REWRITE PLUGINS

New APIs for Query Rewriting

- One for pre-parsing rewriting
- Another for post-parsing rewriting

Example Plugin Installation

```
$ ~/sandboxes/msb_5_7_8/my sql <  
./share/install_rewriter.sql
```

More on this: <http://mysqlserverteam.com/the-query-rewrite-plugins/>

Query Rewriting Setup

```
mysql-5.7.8 > INSERT INTO query_rewrite.rewrite_rules (pattern, replacement) VALUES ('SELECT ?', 'SELECT ? + 1');
Query OK, 1 row affected (0.01 sec)
```

```
mysql-5.7.8 > CALL query_rewrite.flush_rewrite_rules();
Query OK, 0 rows affected (0.01 sec)
```

```
mysql-5.7.8 > SELECT 1;
+-----+
| 1 + 1 |
+-----+
|      2 |
+-----+
1 row in set, 1 warning (0.00 sec)
```

```
mysql-5.7.8 > SHOW WARNINGS;
+-----+-----+-----+-----+
| Level | Code | Message                                                                                               |
+-----+-----+-----+-----+
| Note  | 1105 | Query 'SELECT 1' rewritten to 'SELECT 1 + 1' by a query rewrite plugin |
+-----+-----+-----+-----+
1 row in set (0.00 sec)
```

Considerations

- It cannot correct malformed queries- pattern and replacement must be syntactically correct
- Useful for query optimizations for 3rd party applications
- User is notified that a rewrite happened with a note-level message
- Low overhead (5%); specially for untouched queries
- You can do stupid things like:

```
mysql-5.7.5 ((none)) > INSERT INTO  
query_rewrite.rewrite_rules( pattern, replacement ) VALUES  
( 'SELECT 1', 'DROP TABLE test.test' );  
Query OK, 1 row affected (0.01 sec)
```

Query Optimization: From 0 to 10 (and up to 5.7)

OPTIMIZER HINTS

New functionality since 5.7.7

- MySQL accepts hints for query execution with the syntax `/*+... */:`

```
mysql-5.7.8> EXPLAIN SELECT /*+ NO_ICP(revision)*/ * FROM revision WHERE rev_comment like
'jaime%' AND rev_timestamp > '2008'\G
***** 1. row *****
      id: 1
    select_type: SIMPLE
          table: revision
    partitions: NULL
          type: range
possible_keys:
rev_timestamp,rev_timestamp_rev_page,rev_timestamp_2,rev_comment_rev_timestamp
      key: rev_comment_rev_timestamp
     key_len: 783
        ref: NULL
         rows: 1
    filtered: 50.00
      Extra: Using where
1 row in set, 1 warning (0.00 sec)
```

More info on this feature:

<https://www.percona.com/blog/2015/04/30/optimizer-hints-mysql-5-7-7-missed-manual/>

Syntax

- The syntax is identical to the one Oracle Database uses, but does not deprecate yet the old hint syntax (USE INDEX, STRAIGHT_JOIN, ...)
 - “planned”
- It has some overlap with optimizer_switch
 - Although it has lower granularity (table instead of statement)

Max Query Execution

- Indicate it in milliseconds:

```
mysql-5.7.8 > SELECT /*+ MAX_EXECUTION_TIME(1784)*/ SLEEP(10)\G
***** 1. row *****
SLEEP(2000): 1
1 row in set (1.78 sec)
```

- `max_statement_time` syntax and variable seems to have been removed since 5.7.8

Other hints

- BKA / NO_BKA
- BNL / NO_BNL
- MRR / NO_MRR
- NO_ICP
- NO_RANGE_OPTIMIZATION
- QB_NAME (controls the query block where to apply the hint to)

Query Optimization: From 0 to 10 (and up to 5.7)

SQL MODE CHANGES

203

- More on this: http://www.tocker.ca/2014/01/14/making-strict-sql_mode-the-default.html

Stricter Defaults

```
mysql> CREATE TABLE `test` (  
  `id` int(11) NOT NULL,  
  `type` enum('movie','album','videogame') NOT NULL,  
  `date` datetime NOT NULL,  
  PRIMARY KEY (`id`)  
) ENGINE=InnoDB;
```

```
mysql> INSERT INTO test (type, date) VALUES ('tv show', -1);
```



What happens?

MySQL 5.5/5.6 with default settings

```
mysql> INSERT INTO test (type, date) VALUES ('tv show', -1);
Query OK, 1 row affected, 3 warnings (0.00 sec)
```

```
mysql> SHOW WARNINGS;
```

+-----+		
Level	Code	Message
+-----+		
Warning	1364	Field 'id' doesn't have a default value
Warning	1265	Data truncated for column 'type' at row 1
Warning	1264	Out of range value for column 'date' at row 1
+-----+		
3 rows in set (0.00 sec)		

```
mysql> SELECT * FROM test;
```

+-----+		
id	type	date
+-----+		
0		0000-00-00 00:00:00
+-----+		
1 row in set (0.00 sec)		

MySQL 5.7

```
mysql> INSERT INTO test (type, date) VALUES ('tv show', -1);  
ERROR 1265 (01000): Data truncated for column 'type' at row 1
```

```
mysql> INSERT INTO test (type, date) VALUES ('videogame', -1);  
ERROR 1292 (22007): Incorrect datetime value: '-1' for column 'date' at row 1
```

```
mysql> INSERT INTO test (type, date) VALUES ('movie', now());  
ERROR 1364 (HY000): Field 'id' doesn't have a default value
```

```
mysql> INSERT INTO test (id, type, date) VALUES (1, 'videogame', now());  
Query OK, 1 row affected (0.01 sec)
```

GROUP BY Behavior Changes

```
mysql-5.6.21 (osm) > SELECT way_id, count(*), node_id
                        FROM way_nodes
                        GROUP BY way_id
                        ORDER BY count(*) DESC
                        LIMIT 10;
```

way_id	count(*)	node_id
155339744	1187	1558095871
243986064	1156	2604713337
87136668	1128	1013304944
148812873	852	1618837453
149200774	835	34921158
183618216	826	1940223096
273858696	824	1267549776
261584374	770	2669122104
227880171	704	2240011804
193564006	684	1808872763

10 rows in set (1.24 sec)



Non-deterministic behavior

With ONLY_FULL_GROUP_BY (default in 5.7)

```
mysql-5.7.5 (osm) > SELECT way_id, count(*), node_id  
                      FROM way_nodes  
                      GROUP BY way_id  
                      ORDER BY count(*) DESC  
                      LIMIT 10;
```

```
ERROR 1055 (42000): Expression #3 of SELECT list is not in GROUP  
BY clause and contains nonaggregated column  
'osm.way_nodes.node_id' which is not functionally dependent on  
columns in GROUP BY clause; this is incompatible with  
sql_mode=only_full_group_by
```


Problem with ONLY_FULL_GROUP_BY in MySQL <= 5.6

```
mysql-5.6.21 (osm) > SET SQL_mode='ONLY_FULL_GROUP_BY';  
Query OK, 0 rows affected (0.00 sec)
```

```
mysql-5.6.21 (osm) > SELECT  
                        u.id as `user id`,  
                        u.display_name as `user name`,  
                        count(*) as `# changesets`  
FROM users u  
      IN changesets c  
      ON u.id = c.user_id  
GROUP BY u.id  
ORDER BY count(*) DESC  
LIMIT 10;
```



**Functional
dependenc
y**

```
ERROR 1055 (42000): 'osm.u.display_name' isn't in GROUP BY
```

More on this: <http://rpbouman.blogspot.com/2014/09/mysql-575-group-by-respects-functional.html>

5.7 Aims for SQL99 Compliance

```
mysql-5.7.5 (osm) > SELECT u.id as `user id`, u.display_name as `user name`,
count(*) as `# changesets` FROM users u JOIN changesets c ON u.id =
c.user_id GROUP BY u.id ORDER BY count(*) DESC LIMIT 10;
```

user id	user name	# changesets
31257	Ed Avis	4333
508	Welshie	2696
1016290	Amaroussi	2351
352985	ecatmur	1450
142807	SDavies	1342
736	Steve Chilton	1182
346	Tom Chance	1175
38784	Tom Morris	1165
88718	UrbanRambler	1151
1611	Harry Wood	1091

10 rows in set (0.09 sec)

Backward Compatibility

- Some ORMs and frameworks change the default SQL Mode:
 - Ruby on Rails 4+ sets STRICT_ALL_TABLES
 - Drupal 7+ sets TRADITIONAL
 - Mediawiki will set TRADITIONAL, ONLY_FULL_GROUP_BY
- Other applications do not work in standard-compliance modes:
 - Wordpress used to not work in strict mode (fixed):
<http://www.xaprb.com/blog/2013/03/15/wordpress-and-mysqls-strict-mode/>
 - Cacti does not work with ONLY_FULL_GROUP_BY,NO_ZERO_DATE

Deprecated Modes

- `ERROR_FOR_DIVISION_BY_ZERO`, `NO_ZERO_DATE`, and `NO_ZERO_IN_DATE` are deprecated and do nothing
 - Use `STRICT_TRANS_TABLES` or `STRICT_ALL_TABLES`, which include those modes and produce an error

Query Optimization: From 0 to 10 (and up to 5.7)

GIS IMPROVEMENTS & JSON TYPES

Find the Closest Starbucks



```
mysql-5.7.5 (osm) > SET @lat:=51.49353; SET @lon:=-0.18340;
mysql-5.7.5 (osm) > SELECT n.node_id,
                           n.longitude/100000000 as longitude,
                           n.latitude/100000000 as latitude,
                           sqrt(pow((latitude/100000000 - @lat) * 111257.67, 2)
                                + pow((longitude/100000000 - @lon) * 69450.32, 2))
                           as `distance in metres`
  FROM nodes n
 JOIN node_tags n_t1
    ON n.node_id = n_t1.node_id
 JOIN node_tags n_t2
    ON n.node_id = n_t2.node_id
 WHERE
    n_t1.k = 'amenity' and
    n_t1.v = 'cafe' and
    n_t2.k = 'name' and
    n_t2.v = 'Starbucks'
 ORDER BY `distance in metres` ASC
 LIMIT 1;
```

node_id	longitude	latitude	distance in metres
699693936	-0.1823	51.4945	130.9792513096838

1 row in set (0.20 sec)

This Query is Slow

```
mysql-5.7.5 (osm) > EXPLAIN SELECT n.node_id,  
n.longitude/10000000 as longitude, n.latitude/10000000  
as latitude, sqrt(pow((latitude/10000000-  
@lat)*111257.67, 2) + pow((longitude/10000000-  
@lon)*69450.32, 2)) as `distance in metres` FROM nodes  
n JOIN node_tags n_t1 ON n.node_id = n_t1.node_id JOIN  
node_tags n_t2 ON n.node_id = n_t2.node_id WHERE  
n_t1.k = 'amenity' and n_t1.v = 'cafe' and n_t2.k =  
'name' and n_t2.v = 'Starbucks' ORDER BY `distance in  
metres` ASC LIMIT 1\G
```

```
***** 1. row *****  
      id: 1  
select_type: SIMPLE  
      table: n_t1  
  partitions: NULL  
        type: ALL  
possible_keys: PRIMARY  
          key: NULL  
      key_len: NULL  
         ref: NULL  
        rows: 832040  
  filtered: 0.00  
    Extra: Using where; Using temporary; Using  
file sort
```

```
***** 2. row *****  
      id: 1  
select_type: SIMPLE  
      table: n_t2  
  partitions: NULL  
        type: ref  
possible_keys: PRIMARY  
          key: PRIMARY  
      key_len: 8  
         ref: osm.n_t1.node_id  
        rows: 3  
  filtered: 1.41  
    extra: Using where  
***** 3. row *****  
      id: 1  
select_type: SIMPLE  
      table: n  
  partitions: NULL  
        type: ref  
possible_keys: PRIMARY  
          key: PRIMARY  
      key_len: 8  
         ref: osm.n_t1.node_id  
        rows: 1  
  filtered: 100.00  
    Extra: NULL  
3 rows in set, 1 warning (0.00 sec)
```

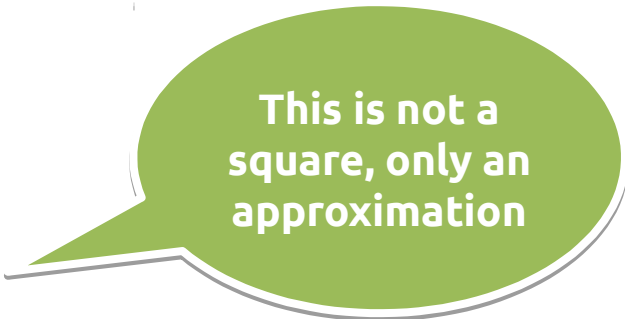
Can We Optimize it?

- We could add a bounding box:

```
mysql-5.7.5 (osm) > ALTER TABLE nodes add index(latitude, longitude);
```

```
mysql-5.7.5 (osm) >
```

```
SELECT n.node_id,  
       n.longitude/100000000 as longitude,  
       n.latitude/100000000 as latitude,  
       sqrt(pow((latitude/100000000 - @lat) * 111257.67, 2) +  
            pow((longitude/100000000 - @lon) * 69450.32, 2))  
       as `distance in metres`  
FROM nodes n  
JOIN node_tags n_t1  
  ON n.node_id = n_t1.node_id  
JOIN node_tags n_t2  
  ON n.node_id = n_t2.node_id  
WHERE  
  n_t1.k = 'amenity' and  
  n_t1.v = 'cafe' and  
  n_t2.k = 'name' and  
  n_t2.v = 'Starbucks' and  
  n.latitude BETWEEN ((@lat - 1000/111257.67) * 100000000)  
                    AND ((@lat + 1000/111257.67) * 100000000) and  
  n.longitude BETWEEN ((@lon - 1000/69450.32) * 100000000)  
                    AND ((@lon + 1000/69450.32) * 100000000)  
ORDER BY `distance in metres` ASC  
LIMIT 1;
```



This is not a
square, only an
approximation

We Create an Index... and Force It

```
mysql-5.7.5 (osm) > EXPLAIN SELECT ...;
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	n_t1	NULL	ALL	PRIMARY	NULL			1	100	Using where;
1	SIMPLE	n	NULL	ref	PRIMARY, latitude	PRIMARY	8	osm.n_t1.node_id	1	5.00	Using temporary; Using filesort
1	SIMPLE	n_t2	NULL	ref	PRIMARY	PRIMARY	8	osm.n_t1.node_id	3	1.41	Using where

MySQL ignores the newly created index, why?

3 rows in set, 1 warning (0.00 sec)

```
mysql-5.7.5 (osm) > EXPLAIN SELECT ... FROM nodes n FORCE INDEX(latitude) ...;
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	n	NULL	range	latitude	latitude	8	NULL	493666	11.11	Using where; Using index; Using filesort
1	SIMPLE	n_t1	NULL	ref	PRIMARY	PRIMARY	8	osm.n.node_id	3	1.41	Using where
1	SIMPLE	n_t2	NULL	ref	PRIMARY	PRIMARY	8	osm.n.node_id	3	1.41	Using where

3 rows in set, 1 warning (0.00 sec)

```
mysql-5.7.5 (osm) > SELECT ... FROM nodes n
```

node_id	longitude	latitude
699693936	-0.1823	51.40

1 row in set (0.09 sec)

Performance improvement is not great

Still many rows are examined

Most of the gain comes from the covering index, not the filtering

Multiple Range Scans Cannot Be Optimized with BTREE Indexes

- We need quadtrees or R-TREE Indexes for indexing in multiple dimensions
 - The later are implemented in MySQL with the name “SPATIAL indexes”, as they only apply to GIS types
- Spatial indexing is available for the first time for InnoDB tables on MySQL 5.7.5

Creating a Spatial Index

```
mysql-5.7.5 (osm) > ALTER TABLE nodes  
                      ADD COLUMN coord GEOMETRY NOT NULL;  
Query OK, 0 rows affected (21.80 sec)  
Records: 0  Duplicates: 0  Warnings: 0
```

```
mysql-5.7.5 (osm) > UPDATE nodes  
                      SET coord = point(longitude/100000000,  
                                         latitude/100000000);  
Query OK, 2865312 rows affected (34.66 sec)  
Rows matched: 2865312  Changed: 2865312  Warnings: 0
```

```
mysql-5.7.5 (osm) > ALTER TABLE nodes add SPATIAL  
index(coord);  
Query OK, 0 rows affected (1 min 50.00 sec)  
Records: 0  Duplicates: 0  Warnings: 0
```



This is new
in 5.7

New Query

```
mysql> SET @area := envelope(linestring(POINT(@lon - 500/69450.32, @lat - 500/111257.67), POINT(@lon + 500/69450.32, @lat + 500/111257.67)));

mysql> SELECT n.node_id,
               x(n.coord) as longitude,
               y(n.coord) as latitude,
               st_distance(POINT(@lon, @lat), coord) as distance
FROM nodes n
JOIN node_tags n_t1
  ON n.node_id = n_t1.node_id
JOIN node_tags n_t2
  ON n.node_id = n_t2.node_id
WHERE
  n_t1.k = 'amenity' and
  n_t1.v = 'cafe' and
  n_t2.k = 'name' and
  n_t2.v = 'Starbucks' and
  st_within(coord, @area)
ORDER BY st_distance(POINT(@lon, @lat), coord) ASC
LIMIT 1;
```



We can use any
shape we want
thanks to 5.6
improvements

Better Performance

mysql-5.7.5 (osm) > SELECT ...;

node_id	longitude	latitude	distance
699693936	-0.1822879	51.4944808	0.0014631428672541478

1 row in set (0.02 sec)

This field is almost useless

mysql-5.7.5 (osm) > EXPLAIN SELECT ...;

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	n	NULL	range	PRIMARY, coord	coord	34	NULL	2	100.00	Using where; Using filesort
1	SIMPLE	n_t1	NULL	ref	PRIMARY	PRIMARY	8	osm.n_node_id	3	1.41	Using where
1	SIMPLE	n_t2	NULL	ref	PRIMARY	PRIMARY	8	osm.n_node_id	3	1.41	Using where

3 rows in set, 1 warning (0.00 sec)

Better Filtering

mysql-5.7.5 (osm) > SHOW STATUS LIKE 'Hand%';

Not using the index: Using the BTREE index: Using the SPATIAL index:

Variable_name	Value
Handler_commit	1
Handler_delete	0
Handler_discover	0
Handler_external_lock	6
Handler_mrr_init	0
Handler_prepare	0
Handler_read_first	1
Handler_read_key	1914
Handler_read_last	0
Handler_read_next	1954
Handler_read_prev	0
Handler_read_rnd	1
Handler_read_rnd_next	833426
Handler_rollback	0
Handler_savepoint	0
Handler_savepoint_rollback	0
Handler_update	0
Handler_write	1

18 rows in set (0.00 sec)

Variable_name	Value
Handler_commit	1
Handler_delete	0
Handler_discover	0
Handler_external_lock	6
Handler_mrr_init	0
Handler_prepare	0
Handler_read_first	0
Handler_read_key	274
Handler_read_last	0
Handler_read_next	246540
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_next	0
Handler_rollback	0
Handler_savepoint	0
Handler_savepoint_rollback	0
Handler_update	0
Handler_write	0

18 rows in set (0.00 sec)

Variable_name	Value
Handler_commit	1
Handler_delete	0
Handler_discover	0
Handler_external_lock	6
Handler_mrr_init	0
Handler_prepare	0
Handler_read_first	0
Handler_read_key	522
Handler_read_last	0
Handler_read_next	5254
Handler_read_prev	0
Handler_read_rnd	259
Handler_read_rnd_next	0
Handler_rollback	0
Handler_savepoint	0
Handler_savepoint_rollback	0
Handler_update	0
Handler_write	0

18 rows in set (0.00 sec)

Geohash Functions

```
mysql-5.7.5 (osm) > SELECT ST_GeoHash(@lon, @lat, 10);
```

```
+-----+  
| ST_GeoHash(@lon, @lat, 10) |  
+-----+  
| gcpugy47w3                |  
+-----+  
1 row in set (0.00 sec)
```

- Useful to index coordinates with a BTREE
 - It could be specially useful combined with indexed STORED columns (emulating quadtrees)

More on Geohashing: <http://mysqlserverteam.com/geohash-functions/>

GeoJSON Functions

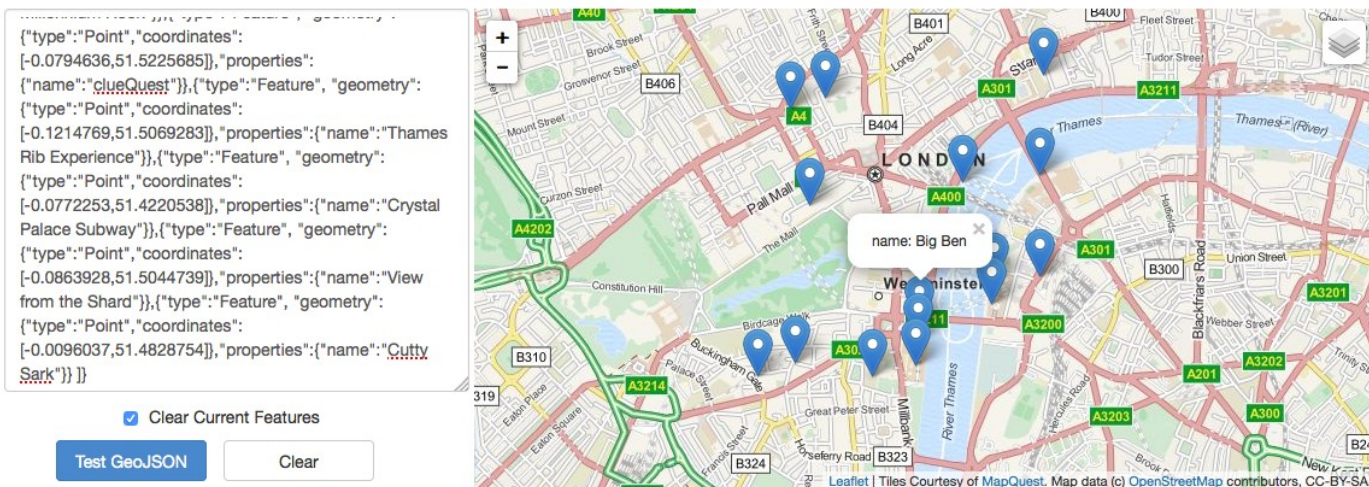
```
mysql-5.7.5 (osm) > SELECT nm.v, ST_AsGeoJson(n.coord)
                        FROM node_tags n_t
                        JOIN nodes n USING (node_id, version)
                        JOIN node_tags nm USING (node_id, version)
                        WHERE n_t.k='tourism' AND
                               n_t.v='attraction' AND
                               nm.k='name';
```

+-----+-----+-----+	
v	ST_AsGeoJson(n.coord)
+-----+-----+-----+	
BedZED	{"type":"Point","coordinates":[-0.1560632,51.3821745]}
Blewcoat School	{"type":"Point","coordinates":[-0.1360484,51.4983179]}
Camden / Buck Street Market	{"type":"Point","coordinates":[-0.143193,51.5400398]}
Camden Lock Village	{"type":"Point","coordinates":[-0.1447181,51.5416552]}
Wimbledon Windmill	{"type":"Point","coordinates":[-0.2315468,51.4376583]}

GeoJSON Functions (cont.)

```
$ ./sandboxes/msb_5_7_5/use osm -B -e "SET SESSION group_concat_max_len = 10000; SELECT
CONCAT('{\"type\": \"FeatureCollection\", \"features\": [ ',
GROUP_CONCAT(CONCAT('{\"type\": \"Feature\", \"geometry\": ', ST_AsGeoJson(n.coord),
', \"properties\": {\"name\": \"', nm.v, '\"}}'), ' ]}') FROM node_tags n_t JOIN nodes n USING
(node_id, version) JOIN node_tags nm USING (node_id, version) WHERE n_t.k='tourism' and
n_t.v='attraction' AND nm.k='name'"
```

<http://geojsonlint.com/>



The screenshot shows the geojsonlint.com interface. On the left, a text area contains a GeoJSON array of features. Below it are buttons for 'Test GeoJSON' and 'Clear', and a checkbox for 'Clear Current Features'. On the right, a map of London is displayed with several blue location pins. A pop-up window for 'Big Ben' is visible on the map.

```
{
  "type": "Point", "coordinates": [
    -0.0794636, 51.5225685
  ], "properties": {
    "name": "ClueQuest"
  }
}, {
  "type": "Point", "coordinates": [
    -0.1214769, 51.5069283
  ], "properties": {
    "name": "Thames Rib Experience"
  }
}, {
  "type": "Point", "coordinates": [
    -0.0772253, 51.4220538
  ], "properties": {
    "name": "Crystal Palace Subway"
  }
}, {
  "type": "Point", "coordinates": [
    -0.0863928, 51.5044739
  ], "properties": {
    "name": "View from the Shard"
  }
}, {
  "type": "Point", "coordinates": [
    -0.0096037, 51.4828754
  ], "properties": {
    "name": "Cutty Sark"
  }
} ]
```

Open Issues

- The SRID can be set and retrieved, but all operations are done in squared euclidean coordinates:

```
mysql-5.7.5 (osm) > SET @p1 := GeomFromText('POINT(-1 51)',  
4326);  
Query OK, 0 rows affected (0.00 sec)
```

```
mysql-5.7.5 (osm) > SET @p2 := GeomFromText('POINT(0 51)', 4326);  
Query OK, 0 rows affected (0.00 sec)
```

```
mysql-5.7.5 (osm) > SET @p3 := GeomFromText('POINT(-1 52)',  
4326);  
Query OK, 0 rows affected (0.00 sec)
```

```
mysql-5.7.5 (osm) > SELECT srid(@p1);  
+-----+  
| srid(@p1) |  
+-----+  
|      4326 |  
+-----+  
1 row in set (0.00 sec)
```

```
mysql-5.7.5 (osm) > SELECT st_distance(@p1, @p2);  
+-----+  
| st_distance(@p1, @p2) |  
+-----+  
|                      1 |  
+-----+  
1 row in set (0.00 sec)
```


```
mysql-5.7.5 (osm) > SELECT st_distance(@p1, @p3);  
+-----+  
| st_distance(@p1, @p3) |  
+-----+  
|                      1 |  
+-----+  
1 row in set (0.00 sec)
```

New JSON Native Data Type

- Since 5.7.8, MySQL allows columns defined with the JSON data type:

```
mysql-5.7.8 > CREATE TABLE json_test(id int PRIMARY KEY  
auto_increment, content JSON);  
Query OK, 0 rows affected (0.03 sec)
```

```
mysql-5.7.8 > INSERT INTO json_test (content) VALUES ('{"type":  
"correct_json"}');  
Query OK, 1 row affected (0.00 sec)
```



They get
validated on
insert

```
mysql-5.7.8 > INSERT INTO json_test (content) VALUES ('{"type":  
"incorrect_json"}');  
ERROR 3140 (22032): Invalid JSON text: "Missing a closing quotation  
mark in string." at position 24 in value (or column) '{"type":  
"incorrect_json"}'.
```

JSON functions

- MySQL includes almost all functions to manipulate JSON that you may think of:
 - Validation test: `JSON_TYPE`
 - Object creation: `JSON_ARRAY`, `JSON_MERGE`, ...
 - Searching: `JSON_EXTRACT`
 - Modifying: `JSON_SET`, `JSON_INSERT`, ...

Indexing JSON

- JSON Columns cannot be indexed:

```
mysql [localhost] {msandbox} (test) > ALTER TABLE  
json_test ADD INDEX(content);  
ERROR 3152 (42000): JSON column 'content' cannot be  
used in key specification.
```
- However, they can be compared with regular fields and use indexes thanks to virtual columns

Query Optimization: From 0 to 10 (and up to 5.7)

CONCLUSIONS

5.7/10.1 About to be released

- Both are currently in Release Candidate
- Unless you are desperate for a feature, skip the first releases (or backport it to your current version)

MySQL 5.7 New Features

- MySQL 5.6 seemed Percona Server-inspired
- MySQL 5.7 seems MariaDB/Galera-inspired
 - Competition is [always good for consumer](#)



Many Optimizer Advantages Have to Be Manually Enabled

- Modifying on a per-query basis:

```
SET optimizer_switch='batched_key_access=on';  
SET join_cache_level=8; # for MariaDB
```

in order to take advantage of them make the features useless unless you are fine-tuning

- I expect that to change in the future

I Herby Declare MyISAM as Dead

- All major MyISAM-only features are now on MySQL 5.7
 - FULLTEXT
 - GIS
 - Transportable tables
- There are still reasons to use MyISAM
 - MyISAM is still required for the mysql schema and non-durable temporary tables (WIP)




Benchmarks

- Do not trust first party benchmarks
 - In fact, do not trust 3rd party benchmarks either
- Only care about the performance of your application running on your hardware



Q&A



Not to Miss

- Official track:
MySQL at Wikipedia: How we do relational data at the Wikimedia Foundation
- Do you want to do query optimization for a website with 20 Billion views per month?
<http://grnh.se/0y4pxm>



Thank You for Attending!

- Do not forget, after the session finishes, to please login with your Percona Live account and “Rate This Session”
- Special thanks to in order by rand() to: Morgan Tocker, Sean Pringle, David Hildebrandt, Bill Karwin, Domas Mituzas, Mark Callaghan, Shlomi Noach, Mark Bergsma, Valerii Kravchuk, Miguel Ángel Nieto, Dimitri Kravtchuk, Olav Sandstå and the whole Wikimedia Team, and all people at the MariaDB, Percona and MySQL/Oracle teams, and the Percona Live Organization and Sponsors